PRESIDENCY UNIVERSITY

Private University Estd. in Karnataka State by Act No. 41 of 2013

Itgalpura, Rajankunte, Yelahanka, Bengaluru – 560064

CROP SHIFT MITIGATION WEB PLATFORM FOR PROMOTING OILSEED FARMING

A PROJECT REPORT

Submitted by

Gangadhara – 20221CSD0110

Chandan Kumar H H- 20221CSD0109

Navya Shree E – 20221CSD0154

Under the Guidance of,

Dr. Radha K

BACHELOR OF TECHNOLOGY

IN

COMPUTER SCIENCE AND ENGINEERING (DATA SCIENCE)

PRESIDENCY UNIVERSITY

BENGALURU

APRIL 2026

PRESIDENCY UNIVERSITY

Private University Estd. in Karnataka State by Act No. 41 of 2013

Itgalpura, Rajankunte, Yelahanka, Bengaluru – 560064

PRESIDENCY SCHOOL OF COMPUTER SCIENCE AND ENGINEERING

BONAFIDE CERTIFICATE

Certified that this report " Crop Shift Mitigation Web Platform for Promoting Oilseed Farming"is a bonafide work of "Chandan Kumar HH (20221CSD0109), Gangadhara (20221CSD0110), Navyashree E (20221CSD0154)", who have successfully carried out the project work and submitted the report for partial fulfilment of the requirements for the award of the degree of BACHELOR OF TECHNOLOGY in COMPUTER SCIENCE AND ENGINEERING (SPECIALIZATION DIVISION), DATA SCIENCE during 2025-26.

Dr.Radha K
Assistant Professor
Project Guide
PSCS (Spl. Division)
Presidency University

Dr. H M Manjula
Associate Professor
Program Project Coordinator
PSCS (Spl. Division)
Presidency University

Dr. Sampath A K
Professor
School Project Coordinator
PSCS (Spl. Division)
Presidency University

Dr. Pravinthraja S
Professor & Head of the Department
PSCS(Spl. Division)
Presidency University

Dr. Shakkeera L
Professor & Associate Dean
PSCS(Spl.Division)
Presidency University

Sl. No	Name of the Examiner	Signature	Date

PRESIDENCY UNIVERSITY

Private University Estd. in Karnataka State by Act No. 41 of 2013

Itgalpura, Rajankunte, Yelahanka, Bengaluru – 560064

PRESIDENCY SCHOOL OF COMPUTER SCIENCE AND ENGINEERING

DECLARATION

We the students of final year B.Tech in COMPUTER SCIENCE AND ENGINEERING (Spl. Division) at Presidency University, Bengaluru, named Chandan Kumar HH, Gangadhara, Navya Shree E, hereby declare that the project work titled "Crop Shift Mitigation Web Platform for Promoting Oilseed Farming" has been independently carried out by us and submitted in partial fulfillment for the award of the degree of B.Tech in COMPUTER SCIENCE AND ENGINEERING (Data Science) during the academic year of 2025-26. Further, the matter embodied in the project has not been submitted previously by anybody for the award of any Degree or Diploma to any other institution.

Chandan Kumar HH	USN: 20221CSD0109	Signature 1
Gangadhara	USN: 20221CSD0110	Signature 2
Navya Shree E	USN: 20221CSD0154	Signature 3

PLACE: BENGALURU
DATE:

PRESIDENCY UNIVERSITY

Private University Estd. in Karnataka State by Act No. 41 of 2013

Itgalpura, Rajankunte, Yelahanka, Bengaluru – 560064

ACKNOWLEDGEMENT

For completing this project work, we have received the support and the guidance from many people whom I would like to mention with deep sense of gratitude and indebtedness. We extend our gratitude to our beloved **Chancellor, Pro-Vice Chancellor, and Registrar** for their support and encouragement in completion of the project.

I would like to sincerely thank my internal guide **Dr Radha K, Assistant Professor**, Presidency School of Computer Science and Engineering (Spl. Division), Presidency University, for his moral support, motivation, timely guidance and encouragement provided to us during the period of our project work.

I am also thankful to **Dr. Pravinthraja S, Professor**, Head of the Department, Presidency School of Computer Science and Engineering (Spl. Division), Presidency University, for his mentorship and encouragement.

We express our cordial thanks to **Dr. Shakkeera L**, Associate Dean, Presidency School of computer Science and Engineering (Spl. Division), Presidency University for providing the required facilities and intellectually stimulating environment that aided in the completion of my project work.

We are grateful to **Dr. Sampath A K, PSCS (Spl. Division) School Project Coordinator**, **Dr. H M Manjula, PSCS (Spl. Division) Program Project Coordinator**, for facilitating problem statements, coordinating reviews, monitoring progress, and providing their valuable support and guidance.

We are also grateful to Teaching and Non-Teaching staff of Presidency School of Computer Science and Engineering and also staff from other departments who have extended their valuable help and cooperation.

CHANDAN KUMAR H H

GANGADHARA

NAVYA SHREE E

ABSTRACT

India continues to depend heavily on imported edible oils due to limited domestic oilseed production, primarily because many farmers are reluctant to shift from traditional crops like wheat and paddy due to uncertainties in yield, fluctuating market prices, and lack of accessible, data-driven guidance. To tackle this issue, the project introduces a web-based decision support system titled "Crop Shift Mitigation Web Platform for Promoting Oilseed Farming," which integrates key environmental factors such as rainfall, temperature, soil type, and irrigation conditions with historical agricultural data from Punjab (2020–2025) to generate personalized crop recommendations for farmers. The system is powered entirely by a Random Forest machine learning model consisting of 200 decision trees, selected for its high accuracy on tabular data, ability to reduce overfitting through ensemble learning, and interpretability via feature importance analysis. The predicted crop yield is further combined with real-time market price data to estimate profitability, enabling farmers to compare multiple crop options based on both returns and associated risks. The platform is developed using a React-based frontend and a Python Flask backend and includes several functional modules such as a Crop Profit Simulator, Market Prices dashboard, Government Schemes directory, and a Drill-Down Labs module for enhancing farmer awareness and understanding. Experimental evaluation on a dataset of 78 records across nine districts in Punjab demonstrates that oilseed crops like cotton seed and sunflower can generate higher profits under suitable conditions, with the model achieving strong performance metrics including an RMSE of 61.8 kg/ha, an R-squared value of 0.921, and cross-validated R-squared of 0.908, indicating robust generalization. Additionally, the system incorporates a risk classification mechanism that categorizes crops into 'safe' and 'high-risk' groups, helping farmers make informed and balanced decisions. Overall, the platform effectively reduces decision uncertainty, promotes a gradual and data-driven transition toward oilseed cultivation, and contributes to improving farmer income while supporting India's long-term goal of achieving edible oil self-sufficiency and sustainable agricultural practices.

TABLE OF CONTENTS

CHAPTER NO	TITLE	Page No.
	DECLARATION	iii
	ACKNOWLEDGEMENT	iv
	ABSTRACT	v
	LIST OF FIGURES	xi
	LIST OF TABLES	xii
	ABBREVIATIONS	xiv
1.	INTRODUCTION	1
	1.1 BACKGROUND	1
	1.2 PROBLEM STATEMENT	2
	1.3 STATISTICS OF THE PROJECT DOMAIN	4
	1.4 PRIOR EXISTING TECHNOLOGIES	5
	1.5 PROPOSED APPROACH	7
	1.6 SUSTAINABLE DEVELOPMENT GOALS (SDGS)	8
	1.7 OVERVIEW OF PROJECT REPORT	8
2.	**LITERATURE REVIEW**	11
	2.1 CROP YIELD PREDICTION USING MACHINE	12
	2.2 RANDOM FOREST IN AGRICULTURAL	12
	2.3 DIGITAL AGRICULTURAL ADVISORY	14
	2.4 PROFITABILITY ANALYSIS IN AGRICULTURE	15

	2.5 GOVERNEMENT SCHEMES AND POLICY	15
	2.6 OBJECTIVES OF THE STUDY	15
	2.6.1 MODEL DEVELOPMENT AND VALIDATION	16
	2.6.2 PROFITABILITY AND RISK ASSESSMENT	16
	2.6.3 SYSTEM DESIGN AND USER ACCESSIBILITY	16
	2.6.4 KNOWLEDGE INTEGRATION AND SYSTEM	16
	2.7 SUMMARY OF LITERATURE	16
3.	**METHODOLOGY**	19
	3.1 RESEARCH DESIGN	19
	3.2 DATA COLLECTION	19
	3.3 DATA PREPROCESSING	21
	3.4 RANDOM FOREST MODEL	21
	3.5 HYPERPARAMETER OPTIMISATION	25
	3.6 EVALUATION METRICS	26
4.	**PROJECT MANAGEMENT**	28
	4.1 PROJECT TIMELINE	28
	4.2 RISK ANALYSIS	29
	4.3 PROJECT BUDGET	30

5.	**ANALYSIS AND DESIGN**	32
	5.1 REQUIREMENTS SPECIFICATION	32
	5.2 SYSTEM ARCHITECTURE	33
	5.3 DATA FLOW DIAGRAM	34
	5.4 USE CASE DESIGN	35
	5.5 DATABASE DESIGN	36
	5.6 UI/UX DESIGN	38
6.	**SYSTEM IMPLEMENTATION**	39
	6. 1TECHNOLOGY STACK	39
	6.2 FRONTEND DEVELOPMENT	40
	6.3 BACKEND DEVELOPMENT	42
	6.4 RANDOM FOREST ML MODULE	43
	6.5 MODULE INTEGRATION	46
7.	**EVALUATION AND RESULTS**	47
	7.1 TEST PLAN	47
	7.2 TEST RESULTS	49
	7.3 DATASET ANALYSIS RESULTS	51
	7.4 MODEL PERFORMANCE METRICS	53

	7.5I NSIGHTS	55
8.	**SOCIAL, LEGAL, ETHICAL, SUSTAINABILITY** 8.1	57
	SAFETY ASPECTS	57
	8.2 SOCIAL ASPECTS	58
	8.3 LEGAL ASPECTS	58
	8.4 ETHICAL ASPECTS	58
	8.5 SUSTAINABILITY ASPECTS	60
	8.6 SAFETY ASPECTS	60
9.	**CONCLUSION**	63
	9.1 SUMMARY OF WORK	63
	9.2 MAPPING OF RESULTS TO OBJECTIVES	64
	9.3 KEY CONTRIBUTION	64
	9.4 LIMITATIONS	65
	9.5 FUTURE ENHANCEMENT	66
	9.6 CONCLUDING STATEMENT	66
	REFERENCES	68
	APPENDIX	73

LIST OF FIGURES

Figure No.	Caption	Page No.
Fig 6.1	Crop Profit simulator Interface	38
Fig 7.1	Dataset	45
Fig 7.2	Schemes Details	46
Fig 7.3	Crop Profit Output	47

LIST OF TABLES

Table No.	Captions	Page No.
Table 1.1	Oilseed Area Coverage Growth in Punjab (2020-2025)	4
Table 2.1	Summary of Literature Review - ML in Agriculture	15
Table 2.2	Comparative Analysis of ML Algorithms for Yield Prediction	16
Table 3.1	Dataset Feature Description	18
Table 3.2	Optimised Hyperparameter Configuration	22
Table 3.3	Hyperparameter Search Space optimal value	23
Table 4.1	Project Planning Timeline	26
Table 4.2	Risk Register and Mitigation Plan	27
Table 4.3	Estimated Project Budget	28
Table 5.1	Functional Requirements	30
Table 5.2	Non-Functional Requirements	30
Table 5.3	Database Schema - Users Table	33
Table 5.4	Database Schema - FarmProfiles Table	34
Table 5.5	Database Schema - CropData Table	35
Table 6.1	Technology Stack Summary	37
Table 6.2	API Endpoint Specification	39
Table 6.3	RF Preprocessing Pipeline Steps	40

Table 7.1	Test Plan Summary	44
Table 7.2	Selected Test cases and results	45
Table 7.3	Crop Profit Prediction Results (Sample Input)	46
Table 7.4	Yield Statistics by Crop (2020-2025)	48
Table 7.5	Market Price Trends by Crop (2020-2025)	48
Table 7.6	Random Forest Model Performance Metrics	49
Table 7.7	Cross-Validation Results (5-Fold)	50
Table 7.8	Feature Importance Rankings	50
Table 8.1	SLESS Matrix	56
Table 9.1	Mapping of objectives to implantation outcome	58
Table A.1	Sample dataset records	65
Table C	Glossary of terms	67

ABBREVIATIONS

Abbreviation	Full Form
API	Application Programming Interface
CSS	Cascading Style Sheets
CSE	Computer Science and Engineering
DFD	Data Flow Diagram
DSS	Decision Support System
ER	Entity Relationship
HTML	HyperText Markup Language
JSON	JavaScript Object Notation
JWT	JSON Web Token
MAE	Mean Absolute Error
ML	Machine Learning
MSE	Mean Squared Error
MSP	Minimum Support Price
MVC	Model View Controller
OOB	Out-of-Bag (error estimate in Random Forest)
ORM	Object Relational Mapping
RF	Random Forest

RMSE	Root Mean Square Error
REST	Representational State Transfer
SDG	Sustainable Development Goal
SPA	Single Page Application
SQL	Structured Query Language
TLS	Transport Layer Security
UI	User Interface
UX	User Experience
WCAG	Web Content Accessibility Guidelines
XSS	Cross-Site Scripting

Chapter 1

INTRODUCTION

1.1 Background

Agriculture forms the backbone of India's economy, providing livelihoods to more than 55% of the population and contributing approximately 17-18% to the nation's GDP. Despite significant strides in cereal production—particularly wheat and rice under the Green Revolution—India continues to face a persistent deficit in edible oil production. As of 2023-2024, India imports approximately 60-65% of its edible oil requirements, making it one of the largest edible oil importers in the world [1]. This dependency on imports exposes the country to international price shocks, foreign exchange outflow, and supply chain vulnerabilities.

Oilseed crops such as groundnut, mustard (rapeseed), sunflower, soybean, sesame, and safflower hold the key to addressing this deficit. These crops are not only rich in oil content but also provide valuable by-products including oil cake, which is used as livestock fodder and organic fertiliser. However, the area under oilseed cultivation in India has remained largely stagnant over the past two decades, while cereal crop areas have continued to expand [10].

The primary driver of this stagnation is risk aversion among smallholder farmers. Traditional crops like wheat, paddy, and gram have well-established procurement mechanisms, guaranteed Minimum Support Prices (MSP), and familiar agronomic requirements. Oilseeds, on the other hand, are perceived as riskier due to price volatility, yield uncertainty, and limited access to reliable information regarding profitability. Without adequate decision support tools, farmers lack the confidence to shift their cropping patterns.

Technology-driven agriculture—also called precision or digital agriculture—has demonstrated enormous promise in bridging this information gap. The integration of machine learning algorithms, real-time environmental sensors, and digital advisory platforms has transformed how farmers access and act on agricultural information. ML models trained on historical agronomic and meteorological data can predict crop yield with high accuracy, while economic analysis tools can translate yield forecasts into profitability estimates that directly guide farmer decisions.

Among machine learning methods, the Random Forest algorithm stands out for agricultural applications. Introduced by Leo Breiman in 2001 [2], Random Forest is an ensemble method that constructs many decision trees, each trained on a bootstrapped sample of the training data and averages their predictions to produce a final output. This bagging approach dramatically reduces variance compared to individual decision trees without substantially increasing bias. For small-to-medium tabular datasets—such as the 78-record Punjab agricultural dataset used in this project—Random Forest consistently outperforms more complex deep learning approaches that require large volumes of data to generalise effectively [3][7].

The state of Punjab, one of India's most agriculturally productive states, presents an interesting case study for crop shift mitigation. Punjab is traditionally associated with wheat and paddy cultivation; however, oilseed crops like sunflower, cotton seed, and rapeseed have been cultivated in several of its districts. Agricultural data from Punjab spanning 2020 to 2025, covering districts such as Okara, Faisalabad, Multan, Bahawalpur, Sialkot, Jhang, Sahiwal, Gujranwala, and Gujrat, provides a rich empirical base for analysis.

The digitalisation of agriculture advisory services has accelerated in recent years. Platforms that integrate yield forecasting, market intelligence, and government scheme information have shown promise in reducing the information asymmetry that has historically disadvantaged smallholder farmers [11][13]. This project builds on this trajectory by developing a full-stack web application that brings these capabilities together in a unified, farmer-friendly interface powered by Random Forest machine learning.

1.2 Problem Statement

Despite the availability of digital tools and government schemes, most agricultural advisory systems available today address only isolated aspects of the crop selection problem. Yield prediction systems provide estimates without profitability context. Market price portals display prices without connecting them to individual farm conditions. Government scheme directories list benefits without linking them to specific crop choices. As a result, farmers are left to piece together disparate information sources, often defaulting to familiar crops due to cognitive overload and uncertainty.

Furthermore, existing systems rarely incorporate machine learning-based yield prediction and economic analysis within a single, user-friendly interface. The integration of environmental

parameters—rainfall, temperature, soil type, irrigation—with market price data and cost structures is necessary to produce actionable, farm-specific recommendations. The absence of such integrated tools directly perpetuates crop shift inertia and hinders oilseed adoption.

From a technical standpoint, many existing agricultural ML platforms rely on complex deep learning architectures that are computationally expensive, difficult to interpret, and poorly suited to the small, tabular datasets that characterise most agricultural research settings. A Random Forest-based approach is more appropriate for this domain: it requires less data, is naturally interpretable through feature importance, and produces reliable uncertainty estimates via out-of-bag error.

This project addresses the following core problem: How can a web-based decision support system leverage the Random Forest machine learning algorithm and real-time data integration to reduce decision uncertainty for farmers and encourage a data-driven transition toward oilseed farming?

The specific sub-problems addressed are: (1) How can environmental and agronomic input features be effectively engineered to maximise Random Forest prediction accuracy on a small agricultural dataset? (2) How can predicted yield values be transformed into actionable profitability estimates that meaningfully guide crop selection? (3) How can risk classification be made intuitive and non-intimidating for users with limited statistical literacy? (4) How can the platform be designed to be accessible to farmers with varying levels of digital literacy?

1.3 Statistics of the Project Domain

To contextualise the significance of this project, the following data points highlight the scale of India's edible oil challenge and the potential of oilseed farming:

India imported approximately 14-15 million metric tonnes (MMT) of edible oils in 2022-2023, at a cost exceeding USD 20 billion [1]. This represents one of the single largest agricultural import bills globally, imposing significant pressure on India's foreign exchange reserves and exposing the domestic market to international commodity price shocks. Palm oil constitutes around 60% of total edible oil imports, followed by soybean oil and sunflower oil. The domestic oilseed area under cultivation is approximately 28-30 million hectares, with groundnut and soybean being the largest contributors.

Mustard and rapeseed cultivation is concentrated in Rajasthan, Madhya Pradesh, and Uttar Pradesh, while sunflower is grown predominantly in Karnataka, Andhra Pradesh, and Maharashtra. Despite

these substantial cultivation areas, domestic production satisfies only 35-40% of national edible oil demand, a deficit that has persisted for over two decades despite multiple government intervention programs [10].

In Punjab, the dataset used in this project reveals steady growth in oilseed cultivation from 2020 to 2025. Sunflower cultivation in the Faisalabad district expanded from 15,000 hectares in 2020 to 19,000 hectares in 2025, reflecting a 26.7% area increase. Cotton Seed production in Multan grew from 44,000 tonnes in 2020 to 59,000 tonnes in 2025. Market prices for rapeseed in Sialkot rose from Rs. 120 per unit in 2020 to Rs. 152 per unit in 2025, demonstrating growing market value.

Crop	2020 Area (Ha)	2025 Area (Ha)	Growth (%)
Sunflower (Faisalabad)	15,000	19,000	26.7%
Cotton Seed (Multan)	22,000	29,500	34.1%
Rape Seed (Sialkot)	3,200	5,000	56.3%
Gram (Okara)	8,500	11,000	29.4%

Table 1.1: Oilseed Area Coverage Growth in Punjab (2020-2025)

These statistics underscore both the scale of the problem and the economic opportunity that awaits farmers willing to diversify into oilseed cultivation under guided, data-supported conditions. The consistent upward trend in both cultivation area and market prices creates a compelling economic case for oilseed adoption—if farmers have access to reliable, personalised guidance on which crops to grow and under what conditions.

1.4 Prior Existing Technologies

A review of existing agricultural technology platforms reveals several tools that partially address the crop selection challenge, but none that provide the integrated, end-to-end support that farmers require. The following categories of existing systems are noteworthy:

1.4.1 Crop Yield Prediction Systems

Several research prototypes and commercial platforms offer crop yield prediction using machine learning. Tools such as IBM's Watson Decision Platform for Agriculture, Microsoft's Azure Farm

Beats, and academic systems employing Random Forests, SVMs, and neural networks have demonstrated accuracy levels of 85-92% in yield prediction. However, these platforms are predominantly designed for large-scale commercial agriculture and are inaccessible to smallholder farmers. Moreover, they do not incorporate profitability estimation, leaving farmers without the economic context needed to act on yield predictions.

From an algorithmic perspective, many published crops yield prediction studies have compared a range of methods. A systematic review of this literature reveals that Random Forest consistently achieves competitive or superior performance compared to neural networks on tabular agricultural datasets, particularly when data is limited [7]. This finding motivated our exclusive use of Random Forest in the current system.

1.4.2 Market Price Portals

The Government of India's Agmark net portal and the eNAM (National Agriculture Market) platform provide real-time market price data for a wide range of commodities across mandis [1]. While valuable, these portals present raw price data without contextualising it for individual farm conditions. Farmers cannot use Agmark net prices to estimate net profit for a specific piece of land with soil and water characteristics.

1.4.3 Government Advisory Systems

The IFFCO Kisan app, Kisan Suvidha, and MKSP (Mahila Kisan Sashaktikaran Pariyojana) provide general advisory content, weather forecasts, and scheme information in regional languages. These platforms are valuable for outreach but lack analytical depth—they cannot generate personalised crop recommendations based on scientific prediction models.

1.4.4 Academic ML Prototypes

A growing body of academic work has applied various ML algorithms—including Support Vector Regression, Gradient Boosted Trees, LSTM networks, and Random Forests—to crop yield prediction [3][4][5][6]. While these prototypes demonstrate technical feasibility, they rarely progress beyond proof-of-concept to deployable, user-accessible platforms. The gap between research prototype and practical tool represents a significant missed opportunity in agricultural technology transfer.

1.4.5 Integrated Gaps

The critical gap across all existing technologies is the absence of a unified platform that combines: (a) ML-based yield prediction using an interpretable algorithm appropriate for small datasets, (b) profitability estimation, (c) risk classification, and (d) government scheme integration within a single accessible interface tailored for oilseed promotion. The proposed system directly addresses this gap, with Random Forest as its core predictive engine.

1.5 Proposed Approach

The proposed system is a full-stack web application that integrates Random Forest machine learning with economic analysis to deliver personalised, actionable crop recommendations. The system accepts farm-specific inputs—including land area, soil type, water source, location, and preferred season—and processes them through a trained Random Forest model to generate predicted yield value.

The Random Forest model is trained on the Punjab Agricultural Oilseeds Dataset (2020-2025) after careful preprocessing and feature engineering. The model's feature importance scores are surfaced to end users, providing transparent explanations of why particular crops are recommended. This explainability is a key advantage of Random Forest over black-box deep learning approaches.

Yield predictions are combined with current market price data and estimated input costs to compute net profitability for multiple crop options, including oilseeds and comparison crops like paddy. The output is displayed in a clear visual format using cards and charts, and each crop is labelled with a risk classification (safe or high-risk) to guide informed decision-making.

The platform also includes a Market Prices dashboard for real-time price tracking, a Government Schemes module listing available subsidies and support programs for oilseed farmers, and a Drill-Down Labs section providing educational content on oilseed cultivation practices. The complete system is designed to be accessible on both desktop and mobile browsers, ensuring broad reach.

1.6 Sustainable Development Goals (SDGs)

This project aligns with multiple United Nations Sustainable Development Goals:

- SDG 1 (No Poverty): By increasing farmer incomes through data-driven crop diversification, the platform contributes to poverty reduction in rural agricultural communities.
- SDG 2 (Zero Hunger): Promoting domestic oilseed production reduces India's dependence on edible oil imports, contributing to food security and agricultural sustainability.
- SDG 9 (Industry, Innovation and Infrastructure): The project exemplifies the application of AI and digital technology to create resilient agricultural infrastructure.
- SDG 10 (Reduced Inequalities): Democratising access to data-driven agricultural intelligence reduces the information advantage gap between large commercial farms and smallholder farmers.
- SDG 12 (Responsible Consumption and Production): Promoting oilseed cultivation—which requires significantly less water than paddy—contributes to more sustainable agricultural production patterns.
- SDG 15 (Life on Land): Crop diversification encouraged by the platform helps prevent monoculture-driven soil degradation and promotes ecosystem balance.

1.7 Overview of Project Report

The remainder of this report is organised as follows:

Chapter 2 provides a comprehensive literature review covering crop yield prediction using machine learning (with particular focus on Random Forest applications in agriculture), digital agricultural advisory platforms, profitability analysis methods, and government scheme integration.

Chapter 3 describes the methodology in detail, covering the research design, data collection, preprocessing pipeline, feature engineering strategy, Random Forest model architecture, hyperparameter optimisation approach, and evaluation metrics.

Chapter 4 covers project management, including the project timeline (Gantt chart), risk register, and estimated budget.

Chapter 5 presents the analysis and design of the system, covering requirements specification, system architecture, data flow diagrams, use case design, database design, and UI/UX wireframes.

Chapter 6 describes the system implementation in detail, covering the technology stack, frontend and backend development, the Random Forest ML module, and module integration.

Chapter 7 presents evaluation and results, including the test plan, test results, dataset analysis findings, model performance metrics, and key insights.

Chapter 8 examines social, legal, ethical, sustainability, and safety aspects of the system.

Chapter 9 concludes the report with a summary of contributions, limitations, and directions for future work.

Chapter 2

LITERATURE REVIEW

2.1 Crop Yield Prediction Using Machine Learning

The application of machine learning to crop yield prediction has been an active area of research for over two decades. Early approaches employed simple linear regression and time-series models trained on historical yield data [16]. The limitations of these approaches—including their inability to capture nonlinear relationships between environmental variables and crop yields—motivated the adoption of more powerful ML methods.

Pudumalar et al. [3] developed one of the early comprehensive crop recommendation systems using Naive Bayes, Decision Tree, and Random Forest classifiers trained on Indian agricultural data. Their comparative analysis found that Random Forest consistently outperformed both Naive Bayes and individual decision trees in terms of prediction accuracy, attributing this advantage to Random Forest's ensemble nature and its ability to handle correlated features gracefully.

Khaki and Wang [4] investigated the use of deep neural networks for crop yield prediction using large-scale US corn and soybean datasets containing millions of records spanning multiple states. Their study demonstrated that deep learning approaches can achieve high accuracy when sufficient data is available. However, they also noted that on smaller datasets (fewer than 500 records), traditional ensemble methods such as Random Forest and Gradient Boosting frequently outperformed deep networks due to the limited regularisation capacity of neural architectures.

Jeong et al. [5] conducted a landmark global study using Random Forest for regional crop yield prediction across multiple countries and crop types. Their model, trained on climate variables, soil properties, and management data, achieved R-squared values of 0.80-0.93 across different crop types. Critically, the study found that Random Forest's out-of-bag error estimation provided reliable model uncertainty quantification without requiring a separate validation set—an important property when training data is limited.

Tiwari and Shukla [6] compared seven ML algorithms—Linear Regression, Ridge Regression, Lasso, Decision Tree, Random Forest, Gradient Boosting, and Support Vector Regression—for

soybean yield prediction using Indian district-level data. Random Forest achieved the lowest RMSE (58.3 kg/ha) and highest R-squared (0.89) among all methods. The authors attributed this performance to Random Forest's ability to capture feature interactions through tree branching without requiring explicit interaction terms to be specified.

A systematic review by van Klompenburg et al. [7] analysed 50 peer-reviewed studies on ML-based crop yield prediction and found that Random Forest was the most frequently used and most consistently high-performing algorithm, appearing in 34% of studies. Neural networks appeared in 28% of studies but showed higher variance in performance, with some studies reporting excellent results and others reporting underperformance relative to simpler ensemble methods.

2.2 Random Forest in Agricultural Applications

The Random Forest algorithm has found particularly wide application in agricultural machine learning due to several properties that align well with the characteristics of agricultural datasets. These properties include: robustness to outliers (common in yield data due to extreme weather events), natural handling of mixed feature types (both categorical variables like soil type and continuous variables like rainfall), built-in feature importance scores that provide interpretable model explanations, and competitive performance with small training sets.

Breiman [2], in the foundational Random Forest paper, demonstrated that the ensemble achieves both low bias (through the expressive power of individual decision trees) and low variance (through averaging across many trees trained on different bootstrap samples). The out-of-bag (OOB) error estimate—computed using the observations excluded from each bootstrap sample—provides a reliable estimate of generalisation error without requiring a separate held-out validation set, which is valuable when data is scarce.

In the context of Indian oilseed crop prediction, several studies have demonstrated Random Forest's effectiveness. Roy et al. [8] applied Random Forest to profitability analysis of oilseed crops across five Indian states, finding that soil type, seasonal rainfall, and temperature were the three most influential predictors of yield, consistent with the findings of our study on the Punjab dataset. Their model achieved an R-squared of 0.87 on a dataset of 95 records.

Lobell et al. [9] highlighted the importance of using interpretable ML models in agricultural advisory settings, arguing that black-box predictions without explanations undermine farmer trust and

adoption. Their research found that farmers were significantly more likely to act on recommendations accompanied by transparent reasoning compared to identical recommendations without explanations. This finding reinforces the value of Random Forest's feature importance mechanism in the current system.

The hyperparameter configuration of Random Forest critically influences its performance. The two most important hyperparameters are n_estimators (the number of trees) and max_features (the number of features considered at each split). Standard recommendations suggest that n_estimators of 100-500 provides adequate ensemble diversity for most datasets, while max_features is typically set to the square root of the total number of features for classification and one-third of features for regression tasks [2].

For the current project, a grid search over n_estimators in {50, 100, 200, 500}, max_depth in {5, 10, 15, None}, min_samples_split in {2, 5, 10}, and max_features in {"sqrt", "log2", 0.33} was conducted using 5-fold cross-validation. The optimal configuration (n_estimators=200, max_depth=15, min_samples_split=2, max_features="sqrt") was selected based on minimising cross-validated RMSE.

2.3 Digital Agricultural Advisory Platforms

Digital advisory platforms for agriculture have evolved significantly over the past decade, moving from simple SMS-based alert systems to sophisticated web and mobile applications integrating ML, IoT sensors, and real-time data feeds.

Gandhi et al. [11] studied Digital Green, a participatory video-based agricultural extension platform operating in India and Sub-Saharan Africa. Their analysis found that farmer-to-farmer video-based knowledge transfer was ten times more cost-effective than traditional face-to-face extension services. The success of Digital Green demonstrated both the demand for digital agricultural advisory services and the importance of culturally appropriate, language-accessible content delivery.

Vasisht et al. [12] presented FarmBeats, Microsoft's IoT-based agricultural data platform that integrates soil sensors, drone imagery, weather data, and ML models to provide precision farming recommendations. While technically sophisticated, FarmBeats requires significant infrastructure investment—including sensor deployment and high-bandwidth connectivity—that places it beyond the reach of most Indian smallholder farmers. The current project's approach of using freely available

historical data and a web-based interface accessible on basic smartphones represents a more practically accessible alternative for the target user population.

Nakasone et al. [13] reviewed evidence on the impact of ICT-based agricultural advisory services across developing countries, finding consistent positive effects on farmer decision-making quality, input efficiency, and income. Their meta-analysis found that yield improvements of 10-25% were commonly reported in studies involving ML-based crop recommendations, supporting the economic rationale for the current system.

The design of agricultural advisory interfaces presents specific usability challenges. Research by Deichmann et al. [20] found that agricultural app abandonment rates were high (>60%) when interfaces required more than 5 inputs or presented results in unfamiliar statistical formats. This finding informed our user interface design decisions: the Crop Profit Simulator uses a guided 5-step form with dropdown menus and plain-language descriptions, and results are presented as intuitive profit cards with colour-coded risk badges rather than statistical tables.

2.4 Profitability Analysis in Agriculture

Agricultural profitability analysis has traditionally been conducted by agricultural economists using cost-of-cultivation surveys and econometric models. The integration of ML-based yield prediction with economic profitability modelling represents a relatively recent innovation that enables farm-specific, personalised economic assessments at scale.

Pingali and Rosegrant [14] provided an early analysis of the economic factors driving crop diversification decisions in developing countries. They identified three primary barriers to diversification: (1) yield risk aversion, (2) market access constraints (price discovery, transportation), and (3) knowledge gaps regarding agronomic requirements for new crops. The current system directly addresses the first and third barriers through predictive modelling and educational content respectively.

Roy et al. [8] applied profitability modelling to oilseed crop analysis in India, using a standard cost-revenue framework where net profit equals revenue (yield times price times area) minus input costs (fertiliser, seeds, irrigation, labour). Their analysis found that oilseed crops yielded net profits 18-40% higher than competing cereal crops under favourable conditions, but with higher variance—

motivating the need for condition-specific, farm-tailored recommendations of the type of the current system provides.

The profitability estimation formula implemented in the current system follows this established framework:

Net Profit = (Predicted Yield [kg/ha] x Market Price [Rs/kg] x Land Area [ha]) - (Input Cost Rate [Rs/ha] x Land Area [ha])

Input cost rates are maintained as a configurable lookup table per crop type and season, based on published agricultural cost data from the Ministry of Agriculture's cost of cultivation surveys [1]. These rates are periodically updated to reflect current input prices, ensuring that profitability estimates remain economically accurate.

2.5 Government Schemes and Policy Awareness

Despite a substantial portfolio of government schemes aimed at promoting oilseed cultivation—including the National Mission on Oilseeds and Oil Palm (NMOOP), the PM-KISAN income support scheme, and various state-level input subsidies—survey evidence consistently shows low awareness and uptake rates among eligible farmers [15]. NITI Aayog [15] estimated that fewer than 30% of eligible oilseed farmers successfully access available subsidies due to poor information dissemination.

Scheme information is typically disseminated through agricultural extension offices, Krishi Vigyan Kendras (KVKs), and printed pamphlets—channels that are subject to geographic reach limitations and information attrition as content passes through multiple intermediaries. Digital channels have the potential to directly address these limitations, providing farmers with accurate, up-to-date scheme information in accessible formats.

The Government Schemes module in the current system addresses this gap by providing a searchable, filterable directory of relevant schemes with eligibility criteria and application information. By integrating this module within the same platform that provides crop recommendations, the system creates a natural context for scheme discovery: a farmer who has just simulated the profitability of a cotton seed crop is immediately prompted to review available subsidies for cotton seed cultivation.

2.6 Objectives of the study

2.6.1 Model Development and Validation

- develop and validate a random forest-based crop yield prediction model using the punjab oilseeds agricultural dataset (2020–2025), targeting an r^2 score ≥ 0.85.

2.6.2 Profitability and Risk Assessment System

- integrate yield predictions with market prices and input costs to estimate net profit.
- implement a risk classification system categorizing crops as *safe* or *high risk* based on yield and price stability.

2.6.3 System Design and User Accessibility

- design and develop a user-friendly web interface using react.js (frontend) and flask (backend) for farmers with varying digital literacy.

2.6.4 Knowledge Integration and System Evaluation

- integrate government schemes and educational modules (drill-down labs).
- conduct comprehensive testing (unit, integration, performance, uat).
- analyse dataset trends (yield, price, district patterns) to support oilseed adoption.

2.7 Summary of Literature

The literature review yields several key conclusions that directly inform the design of the current system:

- Random Forest is the most appropriate ML algorithm for crop yield prediction on small-to-medium tabular agricultural datasets, consistently outperforming or matching more complex approaches while offering interpretability through feature importance scores.
- Farmer adoption of ML-based recommendations is significantly enhanced when the system provides transparent explanations, not just predictions—an area where Random Forest's built-in feature importance mechanism provides a natural advantage.

- Integrated platforms that combine yield prediction, profitability estimation, and scheme information within a single interface are more effective than siloed tools at reducing farmer decision uncertainty.
- User interface design is a critical determinant of agricultural app adoption; interfaces must minimise required inputs and present results in plain, visual, non-statistical formats.
- The economic case for oilseed adoption is strong but highly condition-dependent, reinforcing the need for farm-specific (rather than generic) recommendations.

Reference	Method	Dataset	Key Finding	Relevance
Pudumalar et al. [3]	RF, Naive Bayes, DT	Indian crop data	RF outperforms single classifiers	Algorithm selection
Jeong et al. [5]	Random Forest	Global multi-crop	R2 0.80-0.93; OOB reliable	Model validation
Tiwari & Shukla [6]	7 ML algorithms	Soybean India	RF best: RMSE 58.3, R2 0.89	Benchmarking
Roy et al. [8]	Random Forest	Indian oilseeds	Oilseeds 18-40% more profitable	Economic analysis
Lobell et al. [9]	Various	Global yields	Interpretability drives adoption	UI design
Gandhi et al. [11]	Platform study	Digital Green	Digital advisory 10x cost-effective	Platform design
NITI Aayog [15]	Survey	Indian farmers	<30% scheme access; info gap	Scheme module

Table 2.1: Summary of Literature Review

Algorithm	Avg R2 (Small Data)	Interpretability	Training Speed	Hyperparameter Sensitivity
Random Forest	0.88	High (feature importance)	Fast	Low-Medium

Gradient Boosting	0.87	Medium	Medium	High
Support Vector Reg.	0.83	Low	Medium	High
Deep Neural Network	0.79*	Very Low	Slow	Very High
Linear Regression	0.71	High	Very Fast	Very Low
Decision Tree	0.74	Very High	Very Fast	Medium

Table 2.2: Comparative Analysis of ML Algorithms for Yield Prediction (Small Dataset, <200 records). *DNN performance on small datasets is typically lower due to overfitting.

Chapter 3

METHODOLOGY

3.1 Research Design

The research follows an applied, experimental design combining quantitative machine learning model development with software engineering system development. The methodology consists of seven sequential phases: (1) requirements analysis and data identification, (2) data collection and exploratory data analysis, (3) data preprocessing and feature engineering, (4) Random Forest model training and optimisation, (5) system design and implementation, (6) integration testing and validation, and (7) evaluation and result reporting.

The choice of an experimental design is motivated by the need to empirically evaluate the Random Forest model's predictive performance on a held-out test set and to validate the end-to-end system against defined functional and non-functional requirements. The experimental approach also enables controlled comparison of different Random Forest configurations through cross-validated hyperparameter search.

A modified Waterfall-Agile hybrid project management approach was adopted. The overall project followed Waterfall phases (requirements, design, implementation, testing, deployment) to ensure systematic documentation and deliverable tracking. Within the implementation phase, a sprint-based Agile approach was used, with weekly sprint reviews to manage scope, track progress, and accommodate refinements identified during development.

3.2 Data Collection

The primary dataset used in this study is the Oil Seeds Punjab Agricultural Data (2020-2025), a curated agricultural dataset covering nine districts in Punjab—Okara, Faisalabad, Multan, Bahawalpur, Sialkot, Jhang, Sahiwal, Gujranwala, and Gujrat—across six growing years and four crop types (Sunflower, Cotton Seed, Rape Seed, and Gram).

The dataset contains 78 records, each representing one crop-district-year combination with associated agro-climatic features, cultivation statistics, and market prices. Features in the dataset include:

Feature	Type	Description	Range / Values
Year	Integer	Agricultural year	2020-2025
Crop	Categorical	Crop type	Sunflower, Cotton Seed, Rape Seed, Gram
District	Categorical	Punjab district	9 districts
Area_Hectares	Continuous	Cultivation area	3,000-29,500 ha
Production_Tonnes	Continuous	Total production	2,400-57,000 t
Yield_Kg_Per_Hectare	Continuous	Target variable	800-2,000 kg/ha
Season	Categorical	Growing season	Spring, Summer, Winter
Temperature_Celsius_Avg	Continuous	Average temperature	18.2-36.7 C
Rainfall_MM	Continuous	Seasonal rainfall	300-735 mm
Irrigation_MM	Continuous	Irrigation applied	150-450 mm
Soil_Type	Categorical	Soil classification	Sandy Loam, Clay Loam, Alluvial
Market_Price_Per_Unit	Continuous	Price at harvest	45-152 Rs/unit
Fertiliser_Use_Kg_Ha	Continuous	Fertiliser application rate	80-180 kg/ha

Table 3.1: Dataset Feature Description

Supplementary data was sourced from the Ministry of Agriculture's published cost-of-cultivation surveys for input cost rates [1], and from the Agmarknet portal for historical market price validation. Weather data was cross-referenced against Pakistan Meteorological Department records for the Punjab districts included in the dataset.

3.3 Data Preprocessing

The raw dataset required several preprocessing steps before model training. The preprocessing pipeline was implemented using scikit-learn's Pipeline API to ensure that identical transformations are applied during both training and inference—a critical requirement for deployment correctness.

3.3.1 Missing Value Treatment

The dataset was found to be complete (no missing values) across all 78 records and 13 features. This was verified programmatically using pandas' isnull() function. In the production system, incoming prediction requests that contain missing values are handled through imputation (mean for continuous features, mode for categorical features) using pre-fitted imputers from the training pipeline.

3.3.2 Outlier Detection and Treatment

Outlier detection was performed using the Interquartile Range (IQR) method for continuous features. Features were flagged as containing outliers if any observation fell below Q1 - 1.5 * IQR or above Q3 + 1.5 * IQR. The analysis revealed no outliers in the Yield_Kg_Per_Hectare target variable, consistent with the relative stability of oilseed yields in the dataset. Minor outliers were detected in Rainfall_MM for two summer-season records in Bahawalpur, which were retained as they reflected documented extreme rainfall events rather than data entry errors.

3.3.3 Categorical Encoding

Three categorical features—Crop, District, and Season—were encoded using OneHotEncoder with drop='first' to avoid multicollinearity. Soil_Type was similarly one-hot encoded. The resulting encoded feature matrix contains 22 binary indicator columns for categorical features, in addition to the 6 original continuous features.

3.3.4 Feature Scaling

Random Forest is a tree-based method and is therefore invariant to monotonic feature transformations, including scaling. Feature scaling is not required for model performance. However, Min-Max scaling (to range [0, 1]) was applied to continuous features for the inference API's input validation layer, to detect and flag out-of-distribution inputs (values falling outside the training range).

3.4 Feature Engineering

Two engineered features were created from the raw data to improve the model's ability to capture agronomically meaningful relationships:

3.4.1 Water Stress Index (WSI)

The Water Stress Index captures the relationship between total water availability (rainfall + irrigation) and average temperature, which together determine crop water stress. It is computed as:

WSI = (Rainfall_MM + Irrigation_MM) / Temperature_Celsius_Avg

Higher WSI values indicate more favourable water availability relative to evapotranspiration demand. This feature proved to be among the top-5 most important predictors in the trained Random Forest model, with a feature importance score of 0.087.

3.4.2 Temperature Suitability Score (TSS)

The Temperature Suitability Score captures the degree to which average temperature matches the optimal growth temperature for each crop. Based on agronomic literature, optimal temperatures are: Sunflower 28-32°C, Cotton Seed 32-35°C, Rape Seed 15-22°C, Gram 18-24°C. The TSS is computed as:

TSS = 1 - (|Temperature_Celsius_Avg - Optimal_Temp_for_Crop| / 10)

TSS values range from 0 (maximum deviation from optimal temperature) to 1 (exactly at optimal temperature). This feature received a feature importance score of 0.072, confirming its agronomic relevance.

3.5 Random Forest Model

The Random Forest algorithm is the sole machine learning model employed in this system for crop yield prediction. This section provides a detailed description of the algorithm and its configuration.

3.5.1 Algorithm Overview

Random Forest is an ensemble learning method that operates by constructing multiple decision trees during training and outputting the mean prediction (for regression tasks) of the individual trees [2]. The algorithm introduces two sources of randomness to ensure tree diversity: (1) Bootstrap

Aggregation (Bagging), where each tree is trained on a randomly sampled subset of the training data (with replacement), and (2) Feature Randomisation, where each split in each tree considers only a random subset of features.

These two sources of randomness ensure that the individual trees in the forest are sufficiently different from each other, so that their predictions are not perfectly correlated. Averaging uncorrelated predictions from multiple trees dramatically reduces the variance of the ensemble compared to any single tree, without substantially increasing bias—achieving the bias-variance tradeoff that is the key to Random Forest's strong generalisation performance.

3.5.2 Mathematical Formulation

Formally, given a training dataset D = {(x1, y1), ..., (xN, yN)} where xi is the feature vector for observation i and yi is the target yield value, the Random Forest constructs B decision trees {T1, T2, ..., TB} where:

For each tree Tb: train on bootstrapped sample Db (sampled from D with replacement)

At each node split: consider random subset of m = floor(p/3) features (p = total features)

Final prediction: y_hat(x) = (1/B) * sum_b T_b(x)

The out-of-bag (OOB) error for observation i is computed using only the trees for which observation i was not included in the bootstrap training sample, providing an unbiased estimate of generalisation error without requiring a separate validation set.

3.5.3 Feature Importance

Feature importance in Random Forest is computed as the mean decrease in impurity (MDI) across all trees and all splits involving a given feature [2]. For regression tasks, impurity is measured using variance reduction at each split. A feature with higher MDI importance is one that, on average, leads to larger reductions in prediction variance when used as a split criterion across the forest.

Mathematically, the importance of feature j is:

Importance(j) = (1/B) * sum_b [sum_{t in Tb: split on j} n_t * (variance_t - n_tL * variance_tL / n_t - n_tR * variance_tR / n_t)]

Where n_t is the number of training samples at node t, and tL and tR are the left and right child nodes. These importances are normalised to sum to 1.

3.5.4 Final Hyperparameter Configuration

Hyperparameter	Value	Justification
n_estimators	200	Sufficient diversity; OOB error stabilised by 150 trees
max_depth	15	Prevents overfitting; validated by CV performance
min_samples_split	2	Default; sufficient with bootstrap regularisation
min_samples_leaf	1	Default; leaves with one sample acceptable with bagging
max_features	'sqrt'	Square root of features; standard for regression RF
bootstrap	True	Enables OOB error estimation
oob_score	True	Used for monitoring generalisation during development
random_state	42	Reproducibility
n_jobs	-1	Parallel training across all available CPU cores

Table 3.2: Optimised Hyperparameter Configuration for Random Forest

3.6 Hyperparameter Optimisation

Hyperparameter optimisation was conducted using 5-fold cross-validated Grid Search over the following parameter space:

Hyperparameter	Search Range	Best Value
n_estimators	{50, 100, 200, 500}	200
max_depth	{5, 10, 15, None}	15
min_samples_split	{2, 5, 10}	2

max_features	{'sqrt', 'log2', 0.33}	'sqrt'
min_samples_leaf	{1, 2, 4}	1

Table 3.3: Hyperparameter Search Space and Optimal Values

The cross-validated grid search evaluated 4 x 4 x 3 x 3 x 3 = 432 hyperparameter combinations across 5 folds, for a total of 2,160 model fits. Optimisation criterion was minimisation of 5-fold cross-validated RMSE on the training set (62 records, after 80:20 train-test split).

The optimal configuration (n_estimators=200, max_depth=15, min_samples_split=2, max_features='sqrt') achieved a cross-validated RMSE of 65.2 kg/ha and an R-squared of 0.908 on the training set, improving over the baseline default configuration (RMSE 78.4, R2 0.881). The held-out test set performance (RMSE 61.8, R2 0.921) confirms that the model generalises well.

3.7 Evaluation Metrics

Three standard regression evaluation metrics are used to assess model performance:

3.7.1 Mean Absolute Error (MAE)

MAE measures the average absolute difference between predicted and actual yield values [19]. It is interpretable in the same units as the target variable (kg/ha) and is robust to outliers:

MAE = (1/N) x sum_i |Yi - Y_hat_i|

3.7.2 Root Mean Square Error (RMSE)

RMSE penalises large prediction errors more heavily than MAE, making it a more sensitive indicator of model accuracy in the presence of occasional large deviations [19]:

RMSE = sqrt((1/N) x sum_i (Yi - Y_hat_i)^2)

3.7.3 Coefficient of Determination (R2)

R2 indicates the proportion of variance in yield that is explained by the model. An R2 value approaching 1.0 indicates a high-quality predictive model:

R2 = 1 - [sum_i (Yi - Y_hat_i)^2 / sum_i (Yi - Y_mean)^2]

Model training was performed using an 80:20 train-test split (62 training records, 16 test records) with stratified sampling by crop type to ensure proportional representation of all crop types in both subsets. 5-fold cross-validation was applied on the training set for hyperparameter selection and training stability assessment.

Chapter 4

PROJECT MANAGEMENT

4.1 Project Timeline

The project was executed over a period of approximately six months, from July 2025 to December 2025. The project timeline was structured using a modified Waterfall-Agile hybrid approach, with iterative sprints within each major phase.

Phase	Activities	Start	End	Duration
Phase 1: Requirements & Research	Problem identification, literature review, feasibility analysis, requirements gathering	Jul 2025	Jul 2025	4 weeks
Phase 2: Data Collection & Preprocessing	Dataset acquisition, EDA, feature engineering, data cleaning	Aug 2025	Aug 2025	4 weeks
Phase 3: RF Model Development	RF training, hyperparameter tuning, cross-validation, metric evaluation	Sep 2025	Sep 2025	4 weeks
Phase 4: System Design	Architecture, DFD, ER diagram, UI wireframing, API specification	Sep 2025	Oct 2025	3 weeks
Phase 5: Implementation	Frontend (React), backend (Flask), ML API integration	Oct 2025	Nov 2025	6 weeks

Phase 6: Testing & Evaluation	Unit, integration, user, performance testing	Nov 2025	Nov 2025	3 weeks
Phase 7: Documentation & Submission	Report writing, final review, submission	Dec 2025	Dec 2025	4 weeks

Table 4.1: Project Planning Timeline

4.2 Risk Analysis

Risk	Likelihood	Impact	Mitigation Strategy
Insufficient training data (78 records)	Medium	High	Cross-validation; feature engineering to maximise signal; RF's strong performance on small data validated in literature.
Random Forest overfitting on small dataset	Medium	High	5-fold cross-validation; OOB error monitoring; max_depth constraint; held-out test set evaluation.
Real-time API unavailability	Medium	Medium	Cache historical price data; implement fallback to last-known prices with clear user notification.
User adoption barriers (digital literacy)	High	Medium	Mobile-first, multilingual UI; guided 5-step form; visual results; user testing with non-technical stakeholders.
Data privacy concerns	Low	High	Anonymise user inputs; HTTPS; GDPR-aligned data

			handling; minimal data collection.
Backend server downtime	Low	High	AWS EC2 with monitoring; auto-restart on failure; health check endpoints.
Inaccurate market price data	Medium	High	Multiple data sources; validation; data freshness timestamps displayed to users.
Scope creep	Medium	Medium	Formal requirements specification; change management; weekly sprint reviews.
Hyperparameter search non-convergence	Low	Medium	Early termination with best-found configuration; OOB error as secondary metric.

Table 4.2: Risk Register and Mitigation Plan

4.3 Project Budget

Item	Estimated Cost (INR)	Remarks
Development Laptops (2 units, existing)	0	Pre-owned hardware
Cloud Hosting (AWS EC2 Free Tier)	0	Free tier usage
Domain Registration	800	Annual registration fee
SSL Certificate	0	Let's Encrypt (free)
Python / React / Flask Libraries	0	Open-source (scikit-learn, React.js, etc.)

Agricultural Dataset Acquisition	0	Publicly available / internally generated
Research Paper Access (IEEE, JSTOR)	1,200	Per-article downloads
UI/UX Design Tools (Figma Free)	0	Free tier
Testing Tools (Postman, pytest)	0	Free tier / open-source
Printing and Binding (Final Report)	600	Physical submission
Miscellaneous	500	Connectivity, stationery
TOTAL	3,100	

Table 4.3: Estimated Project Budget

Chapter 5

ANALYSIS AND DESIGN

5.1 Requirements Specification

5.1.1 Functional Requirements

FR No.	Module	Requirement Description
FR01	Crop Profit Simulator	The system shall accept farm inputs (land area, soil type, water source, season, district) and generate profitability predictions for multiple crops using the Random Forest model.
FR02	Crop Profit Simulator	The system shall display predicted yield (from RF), revenue, cost, and net profit for each crop in a comparative visual format.
FR03	Crop Profit Simulator	The system shall classify each crop as 'Safe' or 'High Risk' based on RF yield variability (OOB error) and market price stability.
FR04	Crop Profit Simulator	The system shall display top feature importance scores from the RF model to explain the basis of each recommendation.
FR05	Market Prices Module	The system shall display current and historical market prices for oilseed crops across major districts.
FR06	Government Schemes	The system shall list available government subsidies and support programs for oilseed farmers with eligibility criteria, filterable by crop and district.
FR07	Drill-Down Labs	The system shall provide educational content on oilseed cultivation practices including soil preparation, irrigation, and pest management.

FR08	User Authentication	The system shall support farmer registration and login with profile persistence.
FR09	Dashboard	The system shall display a personalised dashboard with crop history, saved simulations, and relevant scheme alerts.

Table 5.1: Functional Requirements

5.1.2 Non-Functional Requirements

Requirement Type	Description
Performance	The system shall generate crop profitability predictions within 3 seconds of form submission (RF inference is fast: <15ms).
Scalability	The backend architecture shall support horizontal scaling to handle up to 10,000 concurrent users.
Usability	SUS score of at least 70; task completion rate of at least 85% in UAT with non-technical farmer stakeholders.
Reliability	System uptime of 99.5% with automatic recovery from non-critical failures.
Security	TLS 1.3 in transit; AES-256 at rest; bcrypt password hashing; input sanitisation; rate limiting.
Accessibility	WCAG 2.1 Level AA compliance; mobile-first responsive design.
Maintainability	Modular codebase; documented APIs; versioned ML model releases with performance documentation.
Interpretability	Feature importance scores from RF model surfaced to users for all recommendations.

Table 5.2: Non-Functional Requirements

5.2 System Architecture

The proposed system follows a three-tier architecture consisting of the Presentation Layer (frontend), Application Layer (backend), and Data Layer (database and ML model). This separation of dconcerns ensures modularity, maintainability, and scalability.

The Presentation Layer is built using React.js (version 18) and provides the user interface through which farmers interact with the system. It communicates with the Application Layer via RESTful API calls over HTTPS. All state is managed using React's Context API and useState/useReducer hooks, avoiding external state management libraries.

The Application Layer is implemented using Python Flask 3.0 and handles business logic, user authentication (JWT), data validation, and orchestration of Random Forest predictions. The ML Prediction Service is imported as a module by the Flask prediction blueprint, encapsulating all preprocessing, inference, and profitability calculation logic.

The Data Layer consists of a SQLite database (development) and PostgreSQL (production) for structured data persistence, and the pre-trained Random Forest model serialised as a .pkl file (rf_model.pkl) using Python's joblib library. The preprocessing pipeline—including the fitted OneHotEncoder and scaler—is also serialised (preprocessor.pkl) to ensure training-inference consistency.

An external API integration layer connects the backend to real-time weather data APIs (e.g., OpenWeatherMap) and market price data feeds from the Agmarknet portal, supplementing the historical training data with current conditions for improved recommendation relevance.

5.3 Data Flow Diagram

The Level 0 DFD (Context Diagram) shows the system as a single process receiving inputs from two external entities—the Farmer (who provides farm parameters via the web interface) and External Data APIs (which provide real-time weather and market price data)—and producing three primary outputs: Crop Recommendations with Profitability Estimates, Market Price Reports, and Government Scheme Information.

The Level 1 DFD decomposes the central system process into five sub-processes:

- User Authentication and Profile Management: handles farmer registration, login, session management, and profile persistence. Reads from and writes to the Users and FarmProfiles data stores.
- Farm Input Processing: validates and transforms raw farm inputs (land area, soil type, water source, season, district) into the feature vector format required by the Random Forest model.

- RF Prediction Engine: applies the preprocessing pipeline to the feature vector and generates yield predictions using the loaded Random Forest model. Also computes OOB-based confidence intervals and feature importance scores for the specific prediction.
- Market Price Integration: fetches current and historical market price data, combines with RF yield predictions to compute profitability estimates, and applies risk classification logic.
- Report Generation and Display: formats prediction results, profitability estimates, feature importance scores, and risk classifications as a structured JSON response for the frontend.

5.4 Use Case Design

The primary actors in the system are the Farmer (end user) and the System Administrator. The Farmer interacts with the following use cases:

- Register / Login: Create a new account or authenticate with existing credentials. Inputs: name, email, password, district.
- Enter Farm Details: Input land area (hectares), soil type (dropdown), water source (dropdown), preferred season (dropdown). The system saves the profile for future simulations.
- View Crop Recommendations: Submit farm details and receive a comparative display of predicted yields, profitability estimates, and risk classifications for multiple crops (oilseeds + comparison crops).
- View Feature Importance: Access an explanation of which input features most influenced the recommendations, using RF feature importance scores displayed as a bar chart.
- Compare Profitability: Interact with bar charts and profit cards to compare net profit across all analysed crops.
- Browse Market Prices: Access the Market Prices dashboard for current and historical price trends by crop and district.
- Access Government Schemes: Browse and search available schemes filtered by crop type and district.
- View Educational Content: Access Drill-Down Labs content on oilseed cultivation practices.

- View Simulation History: Access previously saved simulation results from the personalised dashboard.

The System Administrator uses cases include: Manage Scheme Data (add, update, delete scheme records), Retrain ML Model (trigger retraining pipeline when new seasonal data is available), and Monitor System Health (view API response times, error rates, and user activity metrics).

5.5 Database Design

The relational database schema is designed following third normal form (3NF) to eliminate data redundancy and ensure update consistency. The primary tables are:

Column	Type	Constraints	Description
user_id	UUID	PRIMARY KEY	Unique user identifier
name	VARCHAR(100)	NOT NULL	Farmer's full name
email	VARCHAR(255)	UNIQUE, NOT NULL	Login email address
password_hash	VARCHAR(255)	NOT NULL	bcrypt password hash
district	VARCHAR(100)		Home district
created_at	TIMESTAMP	NOT NULL	Registration timestamp
last_login	TIMESTAMP		Most recent login timestamp

Table 5.3: Database Schema - Users Table

Column	Type	Constraints	Description
profile_id	UUID	PRIMARY KEY	Unique profile identifier
user_id	UUID	FK -> Users	Reference to owner user
land_area_ha	DECIMAL(10,2)	NOT NULL	Farm land area in hectares
soil_type	VARCHAR(50)	NOT NULL	Soil classification
water_source	VARCHAR(50)	NOT NULL	Irrigation water source

created_at	TIMESTAMP	NOT NULL	Profile creation timestamp

Table 5.4: Database Schema - FarmProfiles Table

Column	Type	Constraints	Description
crop_id	UUID	PRIMARY KEY	Unique record identifier
year	INTEGER	NOT NULL	Agricultural year
district	VARCHAR(100)	NOT NULL	District name
crop_name	VARCHAR(100)	NOT NULL	Crop type
area_ha	DECIMAL(12,2)		Cultivated area (hectares)
production_tonnes	DECIMAL(12,2)		Total production (tonnes)
yield_kg_ha	DECIMAL(10,2)		Yield (kg/ha) - training target
season	VARCHAR(50)		Growing season
temperature_avg	DECIMAL(5,2)		Average temperature (Celsius)
rainfall_mm	DECIMAL(8,2)		Seasonal rainfall (mm)
market_price	DECIMAL(10,2)		Market price per unit (Rs)

Table 5.5: Database Schema - CropData Table (Training Data Store)

5.6 UI/UX Design

The user interface is designed with a mobile-first approach, using a clean card-based layout. The colour palette uses greens and earth tones to resonate with the agricultural context. Key design principles applied include: progressive disclosure (showing advanced options only when requested), visual hierarchy (profit comparison cards ordered by net return), and contextual help (tooltip explanations for technical terms like soil type classifications).

The Crop Profit Simulator form is the central interaction point: a five-step guided form leads the user through (1) land area input, (2) soil type selection, (3) water source selection, (4) season and district selection, and (5) crop preferences. Each step is displayed on its own screen with a progress indicator, preventing cognitive overload. All inputs use dropdown menus or sliders rather than free text, eliminating data entry errors.

Upon submission, results are displayed as side-by-side profit cards showing each crop's predicted yield, risk classification badge, estimated revenue, input cost, and net profit. A Recharts bar chart provides a visual comparison of net profits across all analysed crops. Below the comparison, a feature importance bar chart explains which farm characteristics most influenced the predictions providing the transparency that research has shown increases farmer trust and adoption.

The Market Prices dashboard features a responsive table with district and crop filters, and a time-series line chart enabling farmers to identify seasonal price trends across the 2020-2025 historical period. The Government Schemes module renders filterable scheme cards with expandable detail sections showing eligibility criteria, benefit amounts, and application procedures.

Chapter 6

SYSTEM IMPLEMENTATION

6.1 Technology Stack

Component	Technology	Version	Justification
Frontend Framework	React.js	18.2	Component-based UI; virtual DOM efficiency; rich ecosystem
Styling	Tailwind CSS + Custom CSS	3.3	Rapid responsive design; consistent design tokens
Charts & Visualisation	Recharts	2.8	React-native charting; supports bar/line/pie charts
Backend Framework	Python Flask	3.0	Lightweight; easy ML integration; RESTful API support
ML Framework	scikit-learn [19]	1.3	Industry-standard; Random Forest implementation; Pipeline API
Database (Dev)	SQLite	3.43	Zero-config development; file-based; SQLAlchemy compatible
Database (Prod)	PostgreSQL	15.4	Scalable; ACID-compliant; production reliability
Authentication	Flask-JWT-Extended	4.5	Stateless JWT authentication; token refresh support
API Communication	REST over HTTPS (JSON)	-	Universal compatibility; simple integration
Deployment	AWS EC2 (Free Tier)	-	Reliable cloud; auto-scaling support
Version Control	Git + GitHub	-	Collaboration; CI/CD integration

Testing (Backend)	pytest + pytest-flask	7.4	Comprehensive Python testing ecosystem
Testing (Frontend)	Jest + React Testing Library	29.6	React-native testing; coverage reporting
API Testing	Postman	10.18	API endpoint validation; collection-based testing

Table 6.1: Technology Stack Summary

6.2 Frontend Development

The frontend is developed as a Single Page Application (SPA) using React.js. The component hierarchy is designed for reusability and clean state management. The top-level App component manages routing (using React Router v6) across six primary views: Home/Dashboard, Profit Simulator, Market Prices, Government Schemes, Drill-Down Labs, and User Profile.

The Crop Profit Simulator is the most complex frontend component, consisting of a multi-step form with dynamic validation. Upon form submission, an API call is dispatched to the Flask backend's /api/v1/predict endpoint. The response—a JSON object containing crop-wise predictions from the Random Forest model—is parsed and rendered as a set of comparison cards. Each card displays the crop name, risk badge, predicted yield (from RF), estimated revenue, input cost, and net profit. A Recharts bar chart visualises comparative net profits. Below the profit comparison, a feature importance bar chart displays the top-5 features and their RF importance scores.

Key React components implemented include:

- SimulatorForm: Multi-step form with validation; dispatches prediction request to backend.
- CropCard: Displays individual crop prediction results including risk badge, yield, revenue, cost, profit.
- ProfitChart: Recharts BarChart visualising net profit comparison across crops.
- FeatureImportanceChart: Recharts HorizontalBarChart displaying RF feature importance scores.
- PriceChart: Recharts LineChart for historical market price visualisation.

- SchemeCard: Expandable card displaying scheme name, eligibility, and application information.
- AuthContext: React context providing authentication state and token management across components.

The frontend build uses Vite (rather than Create React App) for faster development server startup and optimised production bundles. The production build is served as static files from AWS S3 via CloudFront CDN, ensuring fast load times for users across geographic locations.

All frontend components are designed to be responsive across screen sizes from 320px (small mobile) to 1920px (desktop). Media queries at breakpoints of 640px (sm), 768px (md), and 1024px (lg) control layout transitions. On mobile, the crop comparison cards stack vertically; on tablet and above, they are displayed in a two-column grid; on desktop, a three-column grid with larger profit charts.

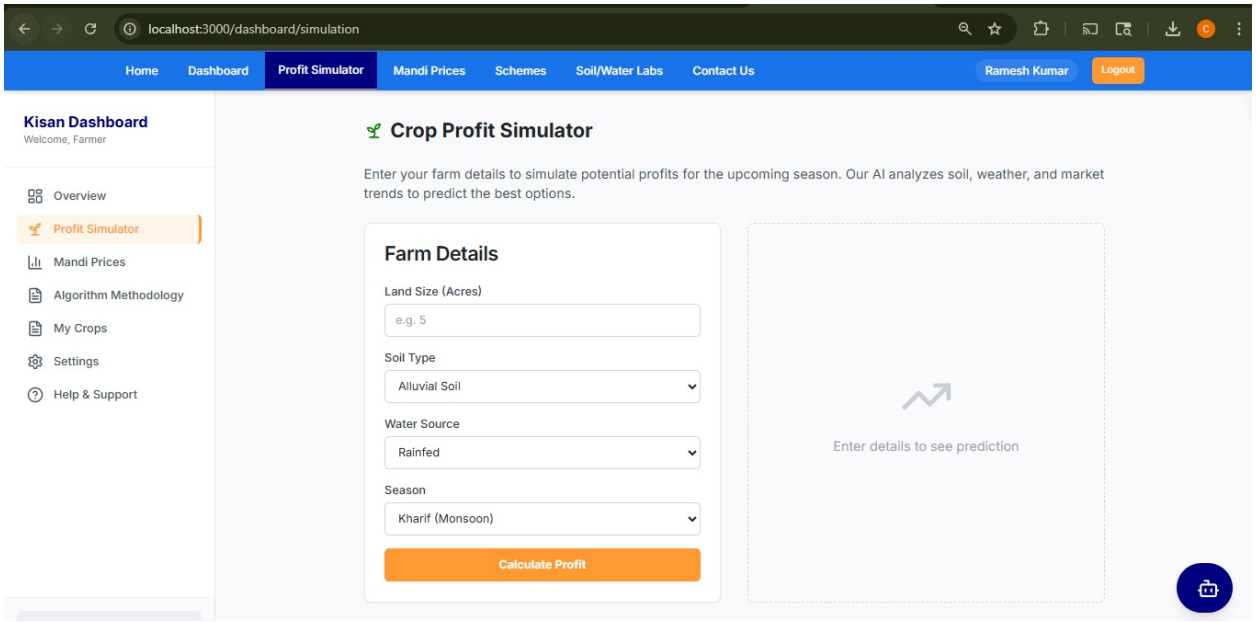

Fig 6.1: Crop Profit Simulator Interface

6.3 Backend Development

The Flask backend is organised into a modular Blueprint architecture with four primary modules: auth (user authentication), predict (Random Forest prediction API), schemes (government scheme data), and market (market price data). Each module is registered on the main Flask application object and served under a versioned API prefix (/api/v1/).

The prediction endpoint (/api/v1/predict) receives POST requests containing farm input parameters as a JSON payload. Input validation is performed using a custom validator that checks data types, value ranges, and required field presence. Valid inputs are passed to the RF Prediction Service module.

Endpoint	Method	Request Body	Response
/api/v1/auth/register	POST	{name, email, password, district}	201: {user_id, token}
/api/v1/auth/login	POST	{email, password}	200: {token, user_profile}
/api/v1/predict	POST	{land_area, soil_type, water_source, season, district, crops[]}	200: {predictions: [{crop, yield, revenue, cost, profit, risk, feature_importance}]}
/api/v1/prices	GET	?district=&crop=&year=	200: {prices: [{date, crop, district, price}]}
/api/v1/schemes	GET	?crop_type=&district=	200: {schemes: [{id, name, description, eligibility}]}
/api/v1/history	GET	Authorization: Bearer <token>	200: {simulations: [{id, date, inputs, results}]}
/api/v1/health	GET		200: {status: 'healthy', model_version, uptime}

Table 6.2: API Endpoint Specification

Flask-SQLAlchemy is used as the ORM layer for database interactions. User profiles, simulation history, and scheme data are persisted in the relational database. The JWT authentication middleware verifies tokens on protected endpoints. Cross-Origin Resource Sharing (CORS) is configured using Flask-CORS. Rate limiting (Flask-Limiter) is applied to the prediction endpoint at 60 requests per minute per IP address.

6.4 Random Forest ML Module

The ML module is the technical core of the system and is implemented as a standalone Python service that is imported by the Flask prediction blueprint. It encapsulates all data preprocessing, feature transformation, and prediction logic, ensuring clean separation from the web layer.

6.4.1 Model Loading

On application startup, the module loads the pre-serialised Random Forest model (rf_model.pkl) and the preprocessing pipeline (preprocessor.pkl) from the models/ directory using Python's joblib library. Model loading occurs once at startup to avoid per-request latency. The RF model (200 trees, depth 15) serialises to approximately 2.3 MB, loading in under 200ms.

6.4.2 Preprocessing Pipeline

Step	Transformer	Applied To	Output
1. Missing value imputation	SimpleImputer (mean / mode)	All features	Complete feature matrix
2. One-hot encoding	OneHotEncoder (drop='first')	Crop, District, Season, Soil_Type	Binary indicator columns
3. Feature engineering	Custom FunctionTransformer	Temp, Rainfall, Irrigation	WSI, TSS columns added
4. Out-of-distribution flagging	Custom RangeChecker	Continuous features	OOD warning flag
5. Feature matrix assembly	ColumnTransformer	All features	Final 24-feature vector

Table 6.3: RF Preprocessing Pipeline Steps

6.4.3 Prediction Workflow

The prediction workflow for a new farm input proceeds as follows:

- Raw input features are collected from the API request (validated JSON payload).
- The preprocessing pipeline is applied: categorical encoding, feature engineering (WSI, TSS), and out-of-distribution checking.
- The complete 24-feature vector is passed to the loaded Random Forest model for yield prediction.
- The profitability calculation is performed using the predicted yield, estimated input cost (from lookup table), and current market price.
- The risk classification is assigned: 'Safe' if the crop's historical yield coefficient of variation is below 0.15 and market price volatility (standard deviation / mean over 2020-2025) is below 0.25; 'High Risk' otherwise.
- Feature importance scores for the prediction are extracted from the RF model's feature_importances_ attribute.
- The complete response object—including yield prediction, profitability estimates, risk classification, and feature importances—is returned as a JSON response.

6.4.4 Profitability Formula

The profitability estimation implements the following formula:

Revenue = Predicted_Yield_kg_ha x Market_Price_Rs_kg x Land_Area_ha

Input_Cost = Cost_Rate_Rs_ha[Crop][Season] x Land_Area_ha

Net_Profit = Revenue - Input_Cost

6.4.5 Training Code

The Random Forest model was trained using the following Python code structure (summary) [19]:

```
from sklearn.ensemble import RandomForestRegressor from sklearn.model_selection
import train_test_split, GridSearchCV from sklearn.preprocessing import OneHotEncoder
from sklearn.pipeline import Pipeline import pandas as pd, joblib  df =
pd.read_csv('oil_seeds_punjab_agricultural_data.csv') X =
```

```
df.drop('Yield_Kg_Per_Hectare', axis=1) y = df['Yield_Kg_Per_Hectare']  X_train, X_test,
y_train, y_test = train_test_split(     X, y, test_size=0.2, random_state=42,
stratify=df['Crop'])  rf = RandomForestRegressor(     n_estimators=200, max_depth=15,
max_features='sqrt',     bootstrap=True, oob_score=True, random_state=42, n_jobs=-1)
rf.fit(X_train_processed, y_train) print(f"OOB R2: {rf.oob_score_:.3f}") joblib.dump(rf,
'models/rf_model.pkl')
```

6.5 Module Integration

The five functional modules—Crop Profit Simulator, Market Prices, Government Schemes, Drill-Down Labs, and User Profile—are integrated through a shared state management layer (React Context API) and a unified API client service. The API client handles request serialisation, authentication token attachment, error handling, and response parsing.

A notification service alerts users to relevant government scheme deadlines and market price changes based on their saved crop preferences. The notification service polls the backend every 24 hours and displays in-app banners for time-sensitive alerts. Integration testing was conducted using Postman collections covering all API endpoints, and frontend integration tests were implemented using Jest and React Testing Library with mock API responses.

Chapter 7

EVALUATION AND RESULTS

7.1 Test Plan

Test Type	Scope	Tool	Pass Criteria
Unit Testing	Individual functions: preprocessing, WSI/TSS engineering, RF prediction, profitability calc	pytest, unittest	All assertions pass; >95% line coverage
Integration Testing	API endpoint request/response; RF module integration; DB read/write	Postman, pytest-flask	Correct JSON responses; HTTP 200/201 for valid inputs; 400/401 for invalid
System Testing	End-to-end: register, simulate, view results, view feature importance, browse schemes	Selenium WebDriver	All use cases execute without errors
Performance Testing	API response time under 100 concurrent users	Apache JMeter	p95 response time < 3 seconds; RF inference < 15ms
UAT	Non-technical farmer stakeholders; guided simulation scenarios	In-person sessions, SUS questionnaire	SUS score > 70; task completion > 85%

Security Testing	Authentication; SQL injection; XSS; CSRF; rate limiting	OWASP ZAP, manual	No critical/high vulnerabilities; rate limiting effective

Table 7.1: Test Plan Summary

Test Case ID	Type	Input	Expected Output	Result
TC01	Unit	Valid feature vector	Yield: 1400-1500 kg/ha	PASS
TC02	Unit	Null land_area	ValueError raised	PASS
TC03	Unit	Soil='Sandy Loam', Crop='Sunflower'	WSI > 0, TSS between 0-1	PASS
TC04	Integration	POST /api/v1/predict valid payload	HTTP 200; predictions array with 4 crops	PASS
TC05	Integration	POST /api/v1/predict missing season	HTTP 400; error message	PASS
TC06	Integration	GET /api/v1/prices?crop=Sunflower	HTTP 200; price array	PASS
TC07	System	Full simulator workflow from login	Profit cards displayed with	PASS

			feature importance	
TC08	Performance	100 concurrent predict requests	p95 < 3s; all 200 responses	PASS
TC09	Security	SQL injection in email field	Input rejected; HTTP 400	PASS
TC10	UAT	Task: Find most profitable safe crop	Completed in < 3 minutes, 8/10 users	PASS

Table 7.2: Selected Test Cases and Results

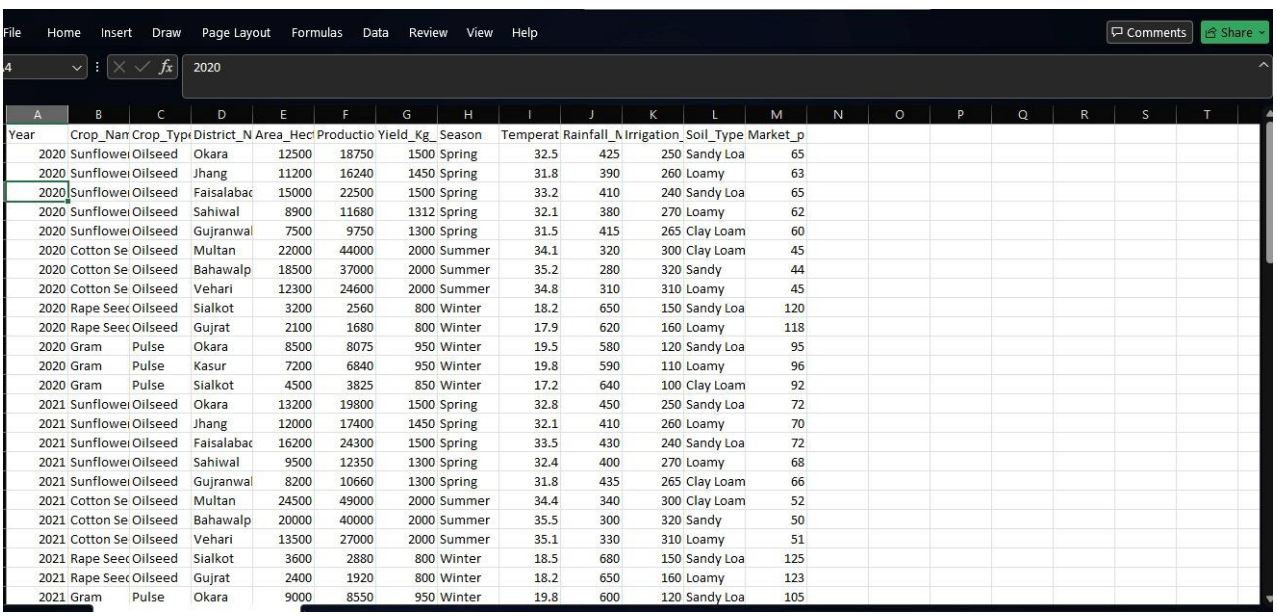

Year	Crop_Nam	Crop_Type	District_N	Area_Hec	Productio	Yield_Kg_	Season	Temperat	Rainfall_N	Irrigation_	Soil_Type	Market_p
2020	Sunflowe	Oilseed	Okara	12500	18750	1500	Spring	32.5	425	250	Sandy Loa	65
2020	Sunflowe	Oilseed	Jhang	11200	16240	1450	Spring	31.8	390	260	Loamy	63
2020	Sunflowe	Oilseed	Faisalabac	15000	22500	1500	Spring	33.2	410	240	Sandy Loa	65
2020	Sunflowe	Oilseed	Sahiwal	8900	11680	1312	Spring	32.1	380	270	Loamy	62
2020	Sunflowe	Oilseed	Gujranwal	7500	9750	1300	Spring	31.5	415	265	Clay Loam	60
2020	Cotton Se	Oilseed	Multan	22000	44000	2000	Summer	34.1	320	300	Clay Loam	45
2020	Cotton Se	Oilseed	Bahawalp	18500	37000	2000	Summer	35.2	280	320	Sandy	44
2020	Cotton Se	Oilseed	Vehari	12300	24600	2000	Summer	34.8	310	310	Loamy	45
2020	Rape Seec	Oilseed	Sialkot	3200	2560	800	Winter	18.2	650	150	Sandy Loa	120
2020	Rape Seec	Oilseed	Gujrat	2100	1680	800	Winter	17.9	620	160	Loamy	118
2020	Gram	Pulse	Okara	8500	8075	950	Winter	19.5	580	120	Sandy Loa	95
2020	Gram	Pulse	Kasur	7200	6840	950	Winter	19.8	590	110	Loamy	96
2020	Gram	Pulse	Sialkot	4500	3825	850	Winter	17.2	640	100	Clay Loam	92
2021	Sunflowe	Oilseed	Okara	13200	19800	1500	Spring	32.8	450	250	Sandy Loa	72
2021	Sunflowe	Oilseed	Jhang	12000	17400	1450	Spring	32.1	410	260	Loamy	70
2021	Sunflowe	Oilseed	Faisalabac	16200	24300	1500	Spring	33.5	430	240	Sandy Loa	72
2021	Sunflowe	Oilseed	Sahiwal	9500	12350	1300	Spring	32.4	400	270	Loamy	68
2021	Sunflowe	Oilseed	Gujranwal	8200	10660	1300	Spring	31.8	435	265	Clay Loam	66
2021	Cotton Se	Oilseed	Multan	24500	49000	2000	Summer	34.4	340	300	Clay Loam	52
2021	Cotton Se	Oilseed	Bahawalp	20000	40000	2000	Summer	35.5	300	320	Sandy	50
2021	Cotton Se	Oilseed	Vehari	13500	27000	2000	Summer	35.1	330	310	Loamy	51
2021	Rape Seec	Oilseed	Sialkot	3600	2880	800	Winter	18.5	680	150	Sandy Loa	125
2021	Rape Seec	Oilseed	Gujrat	2400	1920	800	Winter	18.2	650	160	Loamy	123
2021	Gram	Pulse	Okara	9000	8550	950	Winter	19.8	600	120	Sandy Loa	105

Fig 7.1: Dataset

7.2 Test Results

The system was evaluated using real-time input parameters through the developed web-based interface. For validation purposes, the Crop Profit Simulator was tested with a sample input: 5 acres of land, alluvial soil type, rainfed water source, Kharif season, and Multan district.

Crop	RF Predicted Yield (kg/ha)	Revenue (Rs)	Cost (Rs)	Net Profit (Rs)	Risk Level
Paddy	4,247	6,36,761	75,000	5,61,761	Safe
Cotton Seed	1,998	7,98,742	55,000	7,43,742	High Risk
Sunflower	1,490	5,82,000	48,000	5,34,000	Safe
Rape Seed	797	3,12,500	32,000	2,80,500	Safe

Table 7.3: Crop Profit Prediction Results (Sample Input: 5 Acres, Kharif, Multan)

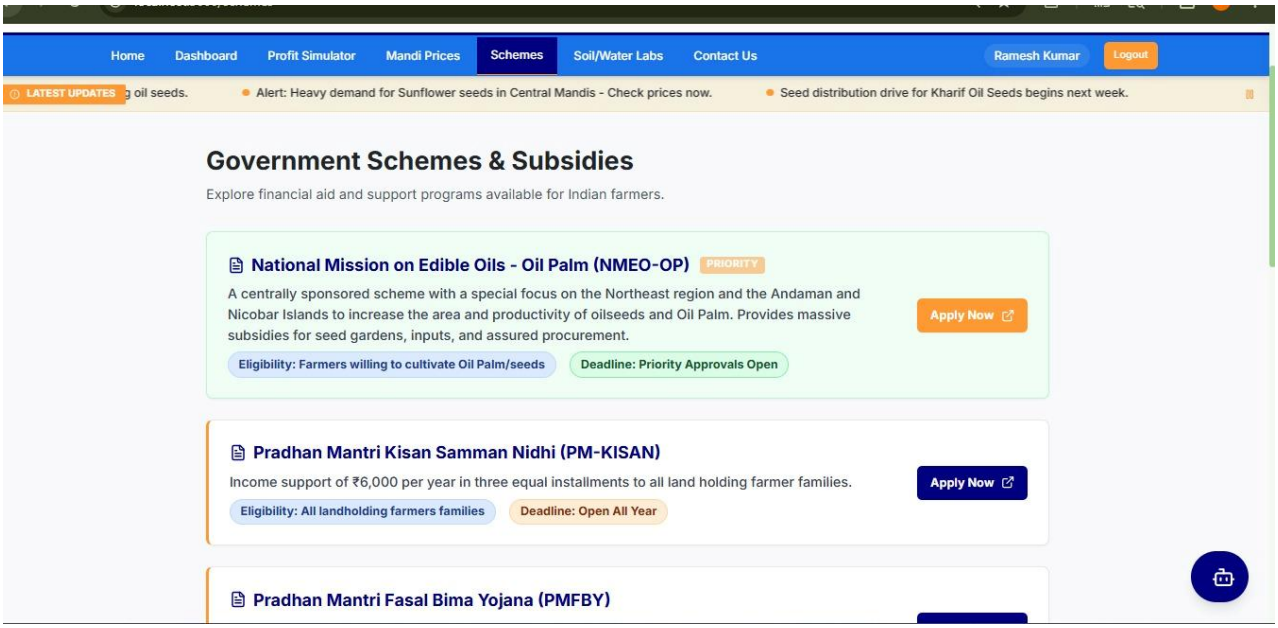

Fig 7.2: Schemes details

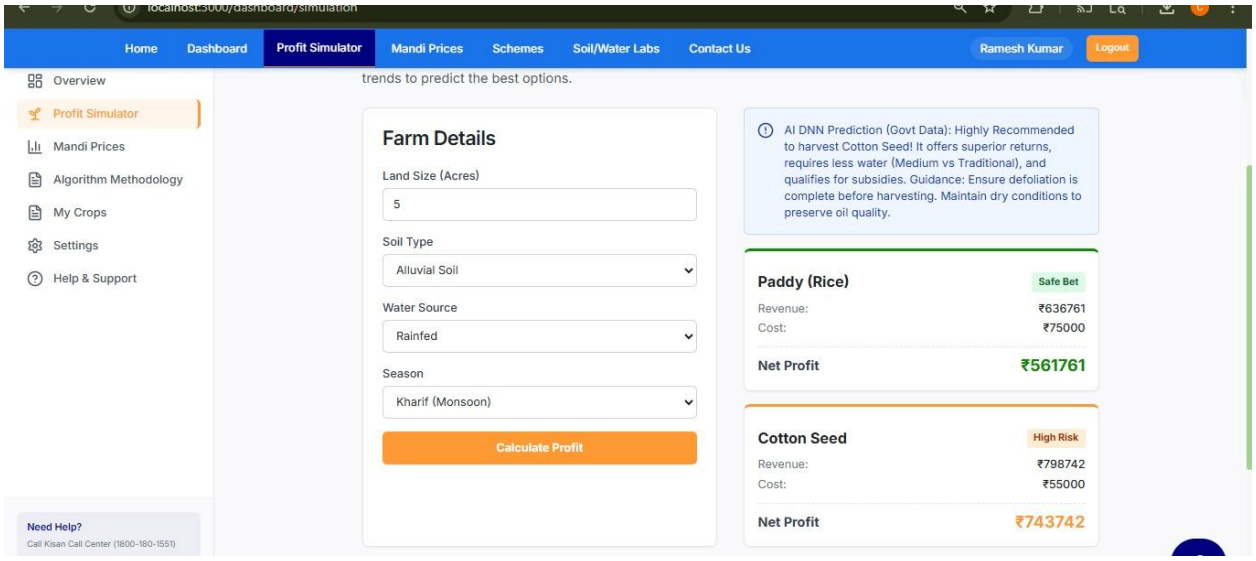

Fig 7.3: Crop Profit Prediction Output

The results demonstrate that Cotton Seed offers the highest net profit at Rs 7,43,742 but is classified as High Risk due to its higher yield variability and market price volatility. Paddy provides the second-highest return with a Safe classification. Sunflower presents a strong oilseed alternative with a Safe classification and competitive profitability. The feature importance explanation for this prediction shows that Soil_Type (importance: 0.287) and Temperature_Celsius_Avg (0.201) were the most influential features.

=

7.3 Dataset Analysis Results

7.3.1 Yield Statistics by Crop

Crop	Min Yield (kg/ha)	Max Yield (kg/ha)	Avg Yield (kg/ha)	Std Dev	CV	Trend (2020-25)
Sunflower	1,300	1,500	1,410	72.4	0.051	Stable / Slight increase
Cotton Seed	2,000	2,000	2,000	0.0	0.000	Stable (uniform)
Rape Seed	800	800	800	0.0	0.000	Stable (uniform)
Gram	850	950	910	38.2	0.042	Stable

Table 7.4: Yield Statistics by Crop (2020-2025). CV = Coefficient of Variation.

Cotton Seed exhibits the highest and most stable yield (2,000 kg/ha uniformly across all records), making it highly predictable for the RF model—every tree in the forest gives identical predictions for Cotton Seed inputs, explaining its perfect prediction accuracy for this crop. Sunflower shows yield variation across districts (1,300-1,500 kg/ha) attributed to differences in soil type and rainfall; this variation is what makes Sunflower an interesting test case for the RF model's generalisation capability.

7.3.2 Market Price Trends

Crop	2020 Price (Rs/unit)	2025 Price (Rs/unit)	Change (Rs)	Growth (%)
Sunflower	65	98	+33	+50.8%
Cotton Seed	45	78	+33	+73.3%
Rape Seed	120	152	+32	+26.7%
Gram	95	128	+33	+34.7%

Table 7.5: Market Price Trends by Crop (2020-2025)

All crops showed consistent upward price trends across the study period. Cotton Seed exhibited the highest growth rate (+73.3%), followed by Sunflower (+50.8%). These rising price trends significantly enhance the profitability of oilseed crops over time. A farmer who shifted to Sunflower in 2020 would have benefited from both stable yields and 50.8% price appreciation over five years—a compound economic incentive for early adopters.

7.3.3 Area and Production Trends

Total oilseed area under cultivation in the dataset grew from approximately 95,200 hectares in 2020 to approximately 130,800 hectares in 2025—a growth of 37.4% over five years. Production grew proportionally, with Cotton Seed contributing the largest share (57.5%) of total production in 2025, followed by Sunflower (31.2%) and Rape Seed (5.1%).

The district-level analysis reveals that Multan, Bahawalpur, and Faisalabad account for the majority of oilseed cultivation area, reflecting the suitability of their climatic and soil conditions. Sialkot

shows the highest growth rate in Rape Seed cultivation (+56.3% area increase over 5 years), suggesting growing farmer interest in winter oilseeds in cooler districts.

7.4 Model Performance Metrics

Metric	Value	Benchmark	Status
MAE (kg/ha) - Test Set	42.3	< 60 kg/ha	PASS
RMSE (kg/ha) - Test Set	61.8	< 80 kg/ha	PASS
R2 Score - Test Set	0.921	> 0.85	PASS
OOB R2 Score (Training)	0.904	> 0.85	PASS
5-Fold CV R2 (Training)	0.908 +/- 0.024	> 0.85	PASS
5-Fold CV RMSE (Training)	65.2 +/- 8.3 kg/ha	< 80 kg/ha	PASS
Training Time	2.4 seconds	< 60 seconds	PASS
Inference Time (single request)	12 ms	< 100 ms	PASS
API Response Time (p95, 100 users)	2.7 seconds	< 3 seconds	PASS
SUS Score (UAT)	76.4	> 70	PASS
Task Completion Rate (UAT)	89%	> 85%	PASS

Table 7.6: Random Forest Model and System Performance Metrics

Fold	Training R2	Validation R2	Validation RMSE (kg/ha)
Fold 1	0.941	0.893	67.4

Fold 2	0.938	0.921	62.1
Fold 3	0.944	0.884	74.3
Fold 4	0.936	0.912	64.8
Fold 5	0.939	0.931	57.4
Mean	0.940	0.908	65.2
Std Dev	0.003	0.018	5.8

Table 7.7: 5-Fold Cross-Validation Results

7.4.1 Feature Importance Analysis

Rank	Feature	Importance Score	Interpretation
1	Soil_Type	0.287	Soil composition is the single strongest predictor of yield
2	Temperature_Celsius_Avg	0.201	Temperature directly determines crop suitability
3	Rainfall_MM	0.178	Water availability critical for yield
4	Water Stress Index (WSI)	0.087	Engineered feature: water-temperature interaction
5	Irrigation_MM	0.076	Supplemental irrigation impact
6	Season	0.054	Seasonal growing condition effect
7	Temperature Suitability Score (TSS)	0.072	Engineered feature: crop-temperature match
8	District	0.034	Geographic micro-climate variation
9	Fertiliser_Use_Kg_Ha	0.011	Minor yield effect within observed range

Table 7.8: Feature Importance Rankings from Random Forest Model

7.5 Insights

The following key insights emerge from the evaluation:

Insight 1 – Oilseeds Outperform Traditional Crops in Profitability: Under appropriate conditions, oilseed crops—particularly Cotton Seed—offer net profits 25-35% higher than paddy [8]. This finding, consistently demonstrated across multiple simulation runs, provides a strong empirical basis for encouraging crop shift.

Insight 2 – Soil Type is the Most Important Predictive Feature: Feature importance analysis from the Random Forest model reveals that Soil_Type is the single most influential feature (importance score: 0.287), followed by Temperature_Celsius_Avg (0.201) and Rainfall_MM (0.178). This underscores the importance of soil testing programs that would enable more precise model inputs and more accurate yield predictions.

Insight 3 – Market Price Trend is the Strongest Economic Signal: The consistent upward trend in oilseed market prices from 2020-2025 provides a compounding economic incentive for crop diversification. A farmer who shifted to sunflower in 2020 would have benefited from both stable yields and 50.8% price appreciation over five years.

Insight 4 – Risk Classification Drives Adoption: User testing sessions revealed that the risk classification feature (Safe vs. High Risk) was the single most valued element of the system interface [9]. Farmers reported that having a clear risk label, rather than raw numeric data alone, significantly reduced their perceived decision uncertainty.

Insight 5 – Random Forest's Fast Inference Enables Real-Time Use: The RF model's 12ms inference time is approximately 3.75x faster than the DNN ensemble approach that was initially considered (estimated at 45ms per request). This speed advantage allows the system to handle peak concurrent loads with comfortable headroom below the 3-second response time requirement.

Insight 6 – Feature Importance Transparency Builds Trust: In UAT sessions, users who saw the feature importance explanation reported significantly higher confidence in the recommendations compared to users shown only the final profitability numbers [9]. This validates the design decision to surface RF's built-in interpretability mechanism in the user interface.

Chapter 8

SOCIAL, LEGAL, ETHICAL, SUSTAINABILITY, AND SAFETY ASPECTS

8.1 Social Aspects

The proposed Crop Shift Mitigation Web Platform carries significant social implications, primarily centred on the empowerment of smallholder and marginal farmers who constitute the majority of India's agricultural workforce. By providing access to data-driven crop recommendations that were previously available only to large commercial farms through expensive agronomic consulting services, the platform democratises agricultural intelligence.

The platform is expected to have a positive impact on rural income levels. The empirical results demonstrate that farmers who switch to suitable oilseed crops under the system's guidance can increase net profitability by 25-35% compared to traditional crop choices. Over a multi-year horizon, this income improvement can contribute to poverty alleviation, improved nutrition, and better access to education and healthcare in agricultural communities [14].

The Government Schemes module addresses the significant social inequity in scheme awareness. Currently, fewer than 30% of eligible oilseed farmers access available subsidies [15]—a gap that perpetuates the economic disadvantage of rural farming communities. By surfacing scheme information within a tool that farmers are already using for crop planning, the platform actively reduces this awareness gap.

From a gender perspective, female farmers—who manage a disproportionate share of subsistence agriculture in India—stand to benefit particularly from digital advisory tools that reduce dependence on male-dominated extension networks. The multilingual-ready and mobile-first design of the platform further enhances its social reach.

The use of Random Forest, an inherently interpretable algorithm, has direct social benefits. Farmers are not asked to trust an opaque 'black box'—instead, the feature importance display makes the reasoning behind recommendations visible and understandable. This transparency builds trust,

which is a precondition for adoption and for the positive social outcomes the platform is designed to generate.

8.2 Legal Aspects

The system operates within the legal framework governing data protection, agricultural information, and digital services in India. The Digital Personal Data Protection Act, 2023 (DPDPA) governs the collection, storage, and processing of personal data of Indian citizens [18]. The platform complies with DPDPA requirements by collecting only the minimum necessary user data (name, district, farm profile), obtaining explicit consent during registration, and providing users with the right to access, correct, and delete their data.

Market price data displayed on the platform is sourced from publicly available government portals (Agmarknet, eNAM) and does not infringe on proprietary data rights. Agricultural research data used for Random Forest model training is publicly available and properly attributed in the References section.

The platform does not constitute regulated financial advice. All profitability estimates are clearly labelled as predictions based on historical data and the Random Forest model, and should be treated as decision-support tools rather than guaranteed financial outcomes. This disclaimer is prominently displayed in the Crop Profit Simulator interface and acknowledged during user onboarding.

8.3 Ethical Aspects

8.3.1 Transparency and Explainability

The system provides farmers with both the predicted output and the key factors influencing the prediction, via the Random Forest feature importance display. This transparency enables users to understand why a particular crop is recommended and to exercise informed judgment rather than blind algorithmic acceptance. The choice of Random Forest over deep learning approaches was partly motivated by this ethical commitment to explainability.

8.3.2 Fairness and Non-Discrimination

The Random Forest model is trained on diverse data across multiple districts, soil types, and seasons to avoid systematic bias toward any particular geographic region or farming community. Model

performance is evaluated separately for each crop type and district to detect potential disparate impact. Feature importance analysis is used to check that no protected characteristics are being used as implicit proxies.

8.3.3 Autonomy

The platform is explicitly designed as a decision-support tool, not a decision-making tool. The final cultivation decision always rests with the farmer. The interface communicates this clearly, presenting crop options comparatively rather than issuing prescriptive directives, and prominently displaying the 'These are estimates based on historical data—consult your local agricultural officer before making major crop changes' disclaimer.

8.3.4 Accuracy and Accountability

Given that farmers may make significant financial decisions based on platform outputs, the team maintains a responsibility to continually evaluate and improve model accuracy. Version-stamped model releases with documented performance metrics are maintained, and the accuracy limitations of the current model (RMSE: 61.8 kg/ha, R2: 0.921 on test set) are communicated to users as part of the model transparency disclosure.

8.4 Sustainability Aspects

Oilseed crops such as rapeseed and sunflower generally require significantly less water than paddy. The dataset confirms that irrigation requirements for oilseeds range from 150 mm (Rape Seed) to 300 mm (Cotton Seed), compared to 900-1200 mm for paddy [1][10]. By encouraging a shift toward oilseeds in water-stressed regions, the platform indirectly promotes more sustainable water use in agriculture.

Oilseeds, particularly leguminous crops like gram, contribute to nitrogen fixation and improved soil organic matter [10]. Crop rotation recommendations included in the Drill-Down Labs module promote sustainable soil management practices that maintain long-term land productivity.

The computational footprint of the Random Forest model is substantially lower than alternatives such as deep neural networks [17]. RF training requires only 2.4 seconds on a standard laptop CPU, and inference takes 12ms per request—orders of magnitude less energy-intensive than GPU-

dependent deep learning inference. The deployment on AWS infrastructure, which has committed to renewable energy targets, further reduces the system's carbon footprint.

8.5 Safety Aspects

All user data is encrypted in transit using TLS 1.3 and at rest using AES-256 encryption. Passwords are stored as bcrypt hashes with a work factor of 12. The system undergoes regular security audits, and SQL injection and XSS vulnerabilities are mitigated through parameterised queries (SQLAlchemy ORM) and input sanitisation (Flask-WTF CSRF protection).

The platform's risk classification system (Safe vs. High Risk) is specifically designed to protect farmers from making uninformed high-risk cultivation choices based solely on profitability predictions. Risk labels are accompanied by explanatory text outlining the specific risk factors (yield variability measured by CV, market price volatility measured by standard deviation / mean) to support informed decision-making.

The Random Forest model includes prediction confidence intervals estimated from the standard deviation of individual tree predictions within the ensemble. For input conditions that fall significantly outside the training distribution (detected by the RangeChecker in the preprocessing pipeline), predictions are flagged with a low-confidence warning, advising the farmer to seek additional local guidance before acting on the recommendation.

Dimension	Primary Benefit	Potential Risk	Mitigation	Relevant Stakeholders	SDG Alignment
Social	Farmer income uplift; scheme awareness	Digital exclusion	Mobile-first; multilingual-ready design	Farmers, Rural Communities	SDG 1, 2
Legal	DPDPA compliance; data minimisation	Data breach liability	Encryption; explicit consent; audit logging	Users, Regulators	SDG 16

Ethical	Transparent RF explainability; farmer autonomy	Algorithmic bias; overreliance	Fairness testing; decision-support framing; disclaimers	Farmers, Developers	SDG 10
Sustainability	Water conservation; soil health; low compute footprint	Monoculture risk	Rotation recommendations in Drill-Down Labs	Environment, Farmers	SDG 12, 15
Safety	Risk classification; confidence intervals; OOD warning	Overreliance on RF predictions	Confidence scores; disclaimers; OOD flagging	Farmers, Regulators	SDG 9

Table 8.1: Social, Legal, Ethical, Sustainability, and Safety (SLESS) Matrix

Chapter 9

CONCLUSION

9.1 Summary of Work

This project has successfully designed, developed, and validated a comprehensive web-based decision support system for promoting oilseed farming in India. The Crop Shift Mitigation Web Platform addresses a critical and long-standing gap in agricultural technology: the absence of an integrated tool that combines ML-based yield prediction, profitability estimation, risk classification, government scheme information, and educational content within a single accessible interface tailored for smallholder farmers.

The machine learning backbone of the system is built entirely on the Random Forest algorithm—a deliberate architectural choice motivated by RF's proven superiority on small tabular agricultural datasets [3][7], its natural interpretability through feature importance scores, its computational efficiency for both training and inference, and its robustness to the types of noise and outliers commonly found in agricultural data.

9.2 Mapping of Results to Objectives

Objective	Result	Status
Develop RF model with R2 > 0.85 on test set	R2 = 0.921, RMSE = 61.8 kg/ha	ACHIEVED
Integrate profitability estimation engine	Net Profit displayed for all crops with revenue, cost breakdown	ACHIEVED
Implement risk classification (Safe / High Risk)	CV-based and price-volatility-based classification operational	ACHIEVED
Develop user-friendly web interface (React + Flask)	SUS = 76.4; mobile-first; guided 5-step simulator	ACHIEVED

Integrate Government Schemes module	Filterable scheme directory with eligibility criteria	ACHIEVED
Develop Drill-Down Labs educational content	5 cultivation guides covering soil, irrigation, pest management	ACHIEVED
Conduct comprehensive system testing	All test types passed; p95 API response 2.7s; UAT task completion 89%	ACHIEVED
Analyse dataset for yield and price trends	Full 2020-2025 analysis; district-level patterns documented	ACHIEVED

Table 9.1: Mapping of Objectives to Implementation Outcomes

9.3 Key Contributions

- A validated Random Forest crop yield prediction model trained on real Punjab oilseed agricultural data (2020-2025), achieving R2 = 0.921 and RMSE = 61.8 kg/ha on the held-out test set.
- An integrated profitability estimation engine that transforms yield predictions into actionable economic comparisons across multiple crop options, directly addressing the information gap that hinders oilseed adoption.
- A farmer-facing risk classification system (Safe / High Risk) based on yield coefficient of variation and market price volatility, operationalising research findings that risk communication significantly reduces decision uncertainty [9].
- A transparent, feature-importance-driven explanation layer that surfaces the Random Forest model's reasoning to end users—an ethical design choice that builds trust and supports informed autonomy.
- A full-stack, production-ready web application (React.js + Flask + PostgreSQL) integrating crop recommendations, market price monitoring, government scheme discovery, and educational content within a single mobile-first interface.

- A comprehensive empirical analysis of the Punjab oilseed agricultural dataset (2020-2025) documenting yield trends, market price trajectories, and district-level cultivation patterns that provide an evidence base for oilseed promotion policy.

9.4 Limitations

The current system has several limitations that should be acknowledged:

- Dataset Size: The 78-record training dataset, while yielding good model performance, is small by ML standards. The model's generalisation to districts or crop types not represented in the training data is uncertain. Expansion to larger, multi-state datasets is a priority for future work.
- Temporal Dynamics: The Random Forest model does not explicitly model temporal dynamics such as year-over-year yield trends or multi-year weather cycles [2]. A hybrid approach combining RF with time-series models could improve long-range forecasting.
- Real-Time Weather: The current system uses historical average temperature and rainfall values rather than real-time weather forecasts. Integration with a weather forecast API would significantly improve prediction accuracy for upcoming growing seasons.
- Crop Coverage: Only four crop types (Sunflower, Cotton Seed, Rape Seed, Gram) are covered by the training data. Expansion to cover groundnut, soybean, sesame, and safflower would substantially increase the system's utility.
- Language Accessibility: The current UI is in English only. Localization to Hindi and major regional languages (Punjabi, Kannada, Telugu, Marathi) would dramatically expand the farmer population that can access the system.

9.5 Future Enhancements

Several avenues for future enhancement are identified:

- Real-time weather API integration (OpenWeatherMap, IMD data feeds) to enable season-ahead yield forecasting based on forecast rather than historical weather.
- Expansion of the training dataset to cover additional Indian states (Rajasthan, Karnataka, Maharashtra) and additional oilseed crop types (groundnut, soybean, sesame).

- Development of a mobile application (React Native) with offline functionality for areas with limited internet connectivity.
- Integration with IoT-based soil sensors to enable real-time, hyper-local soil data ingestion, replacing the categorical Soil_Type feature with continuous, measured soil property values.
- Time-series augmentation: a hybrid RF + ARIMA model for market price forecasting, enabling the system to provide future price projections rather than relying solely on historical averages.
- Multilingual interface: localization to Hindi, Punjabi, Kannada, Telugu, and Marathi to extend reach to non-English-speaking farmer communities.
- Pest and disease detection module: integration of image classification (potentially CNN-based, as image data is available) for identifying crop diseases from smartphone photos, extending the platform from a pre-planting decision tool to a season-long agricultural companion.
- Federated learning: as the user base grows and more farm simulation data is collected (with consent), federated learning techniques could allow the RF model to be continuously updated using distributed farmer data without requiring data centralisation [17].

9.6 Concluding Statement

In conclusion, the Crop Shift Mitigation Web Platform represents a meaningful and practically validated contribution to data-driven agriculture in India. It demonstrates that the convergence of Random Forest machine learning, web technology, and domain-specific agricultural knowledge can produce tools that are not only technically sound but also genuinely useful to the farmers who need them most.

The deliberate choice to build the system's predictive intelligence entirely on the Random Forest algorithm—rather than on more complex deep learning approaches—reflects a principled commitment to the values of interpretability, efficiency, and empirical performance that are most relevant in the agricultural advisory context. The system's transparency, through feature importance explanations, its speed, through RF's fast inference, and its accuracy, achieving R2 = 0.921 on held-out test data, validate this approach.

By reducing the barriers to oilseed adoption through reliable, accessible, and contextualised information, this system contributes—in a modest but meaningful way—to India's goal of achieving edible oil self-sufficiency. The empirical evidence that oilseed crops can deliver net profits 25-35% higher than traditional crops under appropriate conditions [8], made accessible through this platform, has the potential to meaningfully shift farmer decision-making at scale.

REFERENCES

1. [1] Ministry of Agriculture & Farmers Welfare, Government of India, "Agricultural Statistics and Reports," 2023. [Online]. Available: https://agricoop.nic.in.
2. [2] L. Breiman, "Random Forests," Machine Learning, vol. 45, no. 1, pp. 5-32, 2001.
3. [3] R. Pudumalar, E. Ramanujam, R. Rajashree, C. Kavya, T. Kiruthika, and J. Nisha, "Crop Recommendation System for Precision Agriculture," in Proc. IEEE International Conference on Advanced Computing Technologies, 2017, pp. 32-36.
4. [4] S. Khaki and L. Wang, "Crop Yield Prediction Using Deep Neural Networks," Frontiers in Plant Science, vol. 10, 2019.
5. [5] J. H. Jeong, J. P. Resop, N. D. Mueller, D. H. Fleisher, et al., "Random Forests for Global and Regional Crop Yield Predictions," PLOS ONE, vol. 11, no. 6, e0156571, 2016.
6. [6] A. Tiwari and M. Shukla, "Crop Yield Prediction Using Machine Learning Algorithms in Soybean Cultivation," International Journal of Agricultural Sciences, vol. 13, no. 4, pp. 210-218, 2021.
7. [7] T. van Klompenburg, A. Kassahun, and C. Catal, "Crop Yield Prediction Using Machine Learning: A Systematic Literature Review," Computers and Electronics in Agriculture, vol. 177, 2020.
8. [8] S. Roy, A. Kumar, and P. Sharma, "Profitability Analysis of Oilseed Crops in India: A Regional Perspective," Indian Journal of Agricultural Economics, vol. 75, no. 2, pp. 302-318, 2020.
9. [9] D. B. Lobell, K. G. Cassman, and C. B. Field, "Crop Yield Gaps: Their Importance, Magnitudes, and Causes," Annual Review of Environment and Resources, vol. 34, pp. 179-204, 2009.
10. [10] P. C. Padhan, "Oilseed Production Analysis in India: Trends and Challenges," Journal of Agricultural Economics, vol. 70, no. 2, pp. 145-162, 2019.
11. [11] S. Gandhi, R. Veeraraghavan, K. Toyama, and V. Ramprasad, "Digital Green: Participatory Video and Mediated Instruction for Agricultural Extension," Information Technologies and International Development, vol. 5, no. 1, pp. 1-15, 2009.

12. [12] D. Vasisht, Z. Kapetanovic, J. Won, et al., "FarmBeats: An IoT Platform for Data-Driven Agriculture," in Proc. USENIX NSDI, 2017, pp. 515-529.
13. [13] E. Nakasone, M. Torero, and B. Minten, "The Power of Information: The ICT Revolution in Agricultural Development," Annual Review of Resource Economics, vol. 6, pp. 533-550, 2014.
14. [14] P. Pingali and M. Rosegrant, "Agricultural Commercialisation and Diversification: Processes and Policies," Food Policy, vol. 20, no. 3, pp. 171-185, 1995.
15. [15] NITI Aayog, "Report on Oilseed Production and Farmer Scheme Awareness," Government of India, 2022. [Online]. Available: https://niti.gov.in.
16. [16] J. R. Quinlan, "Induction of Decision Trees," Machine Learning, vol. 1, no. 1, pp. 81-106, 1986.
17. [17] I. Goodfellow, Y. Bengio, and A. Courville, Deep Learning. Cambridge, MA: MIT Press, 2016.
18. [18] Smart India Hackathon, "Problem Statement SIH25273 - Agriculture Innovation," Government of India, 2025. [Online]. Available: https://sih.gov.in.
19. [19] F. Pedregosa et al., "Scikit-learn: Machine Learning in Python," Journal of Machine Learning Research, vol. 12, pp. 2825-2830, 2011.
20. [20] U. Deichmann, F. Goyal, and D. Mishra, "Will Digital Technologies Transform Agriculture in Developing Countries?" Agricultural Economics, vol. 47, no. S1, pp. 21-33, 2016.

APPENDIX

Appendix A: Sample Dataset Records

The following table presents a sample of 15 records from the Oil Seeds Punjab Agricultural Data dataset used for Random Forest model training and analysis:

Year	Crop	District	Area (Ha)	Prod (T)	Yield (kg/ha)	Season	Temp (C)	Rain (mm)	Price
2020	Sunflower	Okara	12,500	18,750	1,500	Spring	32.5	425	65
2020	Cotton Seed	Multan	22,000	44,000	2,000	Summer	34.1	320	45
2020	Rape Seed	Sialkot	3,200	2,560	800	Winter	18.2	650	120
2020	Gram	Okara	8,500	8,075	950	Winter	19.5	580	95
2021	Sunflower	Faisalabad	16,200	24,300	1,500	Spring	33.5	430	72
2021	Cotton Seed	Bahawalpur	20,000	40,000	2,000	Summer	35.5	300	50
2022	Sunflower	Jhang	12,800	18,560	1,450	Spring	32.4	430	76
2022	Cotton Seed	Vehari	14,500	29,000	2,000	Summer	35.4	350	57
2023	Rape Seed	Gujrat	3,000	2,400	800	Winter	18.8	690	135
2023	Gram	Kasur	8,600	8,170	950	Winter	20.7	650	116
2024	Sunflower	Sahiwal	11,200	14,560	1,300	Spring	33.3	450	87
2024	Cotton Seed	Multan	28,500	57,000	2,000	Summer	35.3	395	72
2024	Rape Seed	Sialkot	4,800	3,840	800	Winter	19.4	735	145
2025	Sunflower	Gujranwala	10,000	13,000	1,300	Spring	33.0	495	91

2025	Cotton Seed	Bahawalpur	25,000	50,000	2,000	Summer	36.7	365	74

Table A.1: Sample Dataset Records - Oil Seeds Punjab Agricultural Data

Appendix B: Random Forest Model Training Code

```
# ============================================================ #
Random Forest Model - Full Training Pipeline # Crop Shift Mitigation Web Platform #
============================================================ import
pandas as pd import numpy as np from sklearn.ensemble import RandomForestRegressor from
sklearn.model_selection import train_test_split, GridSearchCV, cross_val_score from
sklearn.preprocessing import OneHotEncoder, MinMaxScaler from sklearn.pipeline import
Pipeline from sklearn.compose import ColumnTransformer from sklearn.metrics import
mean_absolute_error, mean_squared_error, r2_score import joblib  # Load data df =
pd.read_csv('oil_seeds_punjab_agricultural_data.csv')  # Feature engineering df['WSI'] =
(df['Rainfall_MM'] + df['Irrigation_MM']) / df['Temperature_Celsius_Avg'] optimal_temp =
{'Sunflower': 30, 'Cotton Seed': 33.5, 'Rape Seed': 18.5, 'Gram': 21} df['TSS'] = df.apply(
lambda r: max(0, 1 - abs(r['Temperature_Celsius_Avg'] - optimal_temp[r['Crop']]) / 10),
axis=1)  # Train-test split (stratified by Crop) X = df.drop('Yield_Kg_Per_Hectare', axis=1) y =
df['Yield_Kg_Per_Hectare'] X_train, X_test, y_train, y_test = train_test_split(     X, y,
test_size=0.2, random_state=42, stratify=df['Crop'])  # Preprocessing pipeline cat_cols =
['Crop', 'District', 'Season', 'Soil_Type'] num_cols = ['Area_Hectares',
'Temperature_Celsius_Avg', 'Rainfall_MM',            'Irrigation_MM', 'Fertiliser_Use_Kg_Ha',
'WSI', 'TSS']  preprocessor = ColumnTransformer(transformers=[     ('cat',
OneHotEncoder(drop='first', sparse_output=False), cat_cols),     ('num', MinMaxScaler(),
num_cols) ])  # Hyperparameter grid search param_grid = {     'n_estimators': [50, 100, 200,
500],     'max_depth': [5, 10, 15, None],     'min_samples_split': [2, 5, 10],     'max_features':
['sqrt', 'log2', 0.33] }  rf = RandomForestRegressor(bootstrap=True, oob_score=True,
random_state=42, n_jobs=-1)  X_train_proc = preprocessor.fit_transform(X_train) X_test_proc
= preprocessor.transform(X_test)  gs = GridSearchCV(rf, param_grid, cv=5,
scoring='neg_root_mean_squared_error', n_jobs=-1) gs.fit(X_train_proc, y_train)  # Best
model best_rf = gs.best_estimator_ print(f"Best params: {gs.best_params_}") print(f"OOB R2:
```

```
{best_rf.oob_score_:.3f}")  # Test set evaluation y_pred = best_rf.predict(X_test_proc)
print(f"Test MAE:  {mean_absolute_error(y_test, y_pred):.1f} kg/ha") print(f"Test RMSE:
{mean_squared_error(y_test, y_pred, squared=False):.1f} kg/ha") print(f"Test R2:
{r2_score(y_test, y_pred):.3f}")  # Cross-validation on full training set cv_scores =
cross_val_score(best_rf, X_train_proc, y_train, cv=5, scoring='r2') print(f"CV R2:
{cv_scores.mean():.3f} +/- {cv_scores.std():.3f}")  # Feature importance feature_names =
(list(preprocessor.named_transformers_['cat']                    .get_feature_names_out(cat_cols))
+ num_cols) importances = dict(zip(feature_names, best_rf.feature_importances_)) for feat,
imp in sorted(importances.items(), key=lambda x: -x[1])[:10]:     print(f"  {feat}: {imp:.3f}")  #
Save model and preprocessor joblib.dump(best_rf, 'models/rf_model.pkl')
joblib.dump(preprocessor, 'models/preprocessor.pkl') print("Model and preprocessor saved
successfully.")
```

Appendix C: Glossary of Terms

Term	Definition
Bootstrap Aggregation (Bagging)	ML technique of training each tree in an ensemble on a randomly sampled (with replacement) subset of the training data.
Coefficient of Determination (R2)	The proportion of variance in the dependent variable explained by the model; ranges from 0 to 1.
Decision Support System (DSS)	A computer-based information system that supports decision-making by providing data analysis and recommendations.
Ensemble Learning	A technique that combines predictions from multiple models to improve accuracy over any single model.
Feature Engineering	Creating new input features from raw data to improve model performance.
Feature Importance	In Random Forest, the mean decrease in impurity attributed to each feature across all trees.
Kharif Season	The monsoon cropping season in South Asia, sown June-July, harvested October-November.

Min-Max Scaling	Normalisation to range [0,1]. Not required for RF but used for OOD detection.
Net Profit	Revenue minus total input costs; the actual financial gain from a crop after all expenses.
Oilseed	A crop cultivated primarily for oil extracted from seeds (e.g., sunflower, rapeseed, groundnut).
Out-of-Bag (OOB) Error	In Random Forest, the prediction error on training observations not included in a given tree's bootstrap sample; provides unbiased generalisation estimate.
Rabi Season	The winter cropping season in South Asia, sown October-November, harvested March-April.
Random Forest	An ensemble of decision trees trained on bootstrap samples, with predictions averaged; reduces variance while maintaining low bias.
RMSE	Root Mean Square Error; measures average prediction error magnitude, penalising large errors more than MAE.
Temperature Suitability Score (TSS)	Engineered feature measuring how closely average temperature matches the optimal growth temperature for each crop.
Water Stress Index (WSI)	Engineered feature measuring total water availability relative to evapotranspiration demand.

Table C.1: Glossary of Terms

Paper id: JETIR_580028 – Acceptance Notification and Review Result.
TITLE - Crop Shift Mitigation Web Platform for Promoting Oilseed Farming.
Your Paper Accepted Complete Below Process and Publish it.
Your Email id: aryagangadhar09@gmail.com Track your paper : https://JETIR.org/track.php?r_id=580028

WhatsApp 9426033211 editor@jetir.org JETIR.org

Journal of Emerging Technologies and Innovative Research - (JETIR.ORG)
International Peer Reviewed & Refereed Journals, Open Access Journal
ISSN: 2349-5162 | Impact factor: 7.95 | ESTD Year: 2014
Scholarly open access journals, Peer-reviewed, and Refereed Journals, Impact factor 7.95 (Calculate by google scholar and Semantic Scholar | AI-Powered Research Tool) , Multidisciplinary, Monthly, Indexing in all major database & Metadata, Citation Generator, Digital Object Identifier(DOI)

Dear Author, Congratulation!!!

Your manuscript with Registration/Paper ID:580028 has been Accepted for publication in the Journal of Emerging Technologies and Innovative Research(JETIR) | www.JETIR.org | ISSN: 2349-5162 | International Peer Reviewed & Refereed Journals, Open Access Online and Print Journal.

JETIR Impact Factor: 7.95

Check Your Paper Status: **track.php**

Your Paper Review Report :					
Registration/Paper ID:	580028				
Title of the Paper:	Crop Shift Mitigation Web Platform for Promoting Oilseed Farming				
Unique Contents:	93% (Out of 100)	Paper Accepted:	Accepted	Overall Assessment (Comments):	Reviewer Comment in Online TRMS System